

194.039 Intelligent Audio and Music Analysis

Author:

Max Tiessler

2025-01-18

Contents

1 Fundamentals	8
1.1 Basics	8
1.1.1 What are the differences between sound, audio, and music?	8
1.1.2 What is the frequency range for audio?	8
1.1.3 Give examples for objective properties in music!	8
1.1.4 What is an audio waveform and how can it be created from sound?	8
1.1.5 What is ADC and what steps does it consist of?	9
1.1.6 What influence do sampling frequency and bit depth for ADC have?	9
1.1.7 What are the requirements for the sampling frequency for ADC, when sampling a signal?	9
1.1.8 What does PCM stand for?	9
1.1.9 Explain how PCM encodes an audio signal!	9
1.2 Audio Features and Signal Framing	10
1.2.1 What are (audio-)features in the context of signal processing?	10
1.2.2 What are examples for basic time-domain audio features?	10
1.2.3 What are problems of global audio features, especially in the context of music? .	11
1.2.4 Explain in your own words what signal framing is, including the relevant parameters!	11
1.2.5 How can signal framing be used to calculate time-local audio features?	11
1.2.6 What is the advantage of time-local audio features?	12
1.2.7 Give examples of time-local audio features and explain what information they reveal!	12
1.2.8 Explain the process of calculating a time-local feature like ZCR!	12
1.3 Frequency Domain Features	13
1.3.1 How does audio signal frequency analysis relate to the human hearing process? .	13
1.3.2 Explain, in your own words, what the Fourier transform is!	13
1.3.3 Sketch a sine-frequency with 10Hz and the magnitude spectrum of this signal! .	13
1.3.4 Sketch a signal that is a mixture of two sines and its magnitude spectrum! . .	13
1.3.5 Explain amplitude, frequency, and phase for a sinusoid!	13
1.3.6 What is a “frequency-domain representation” of an audio signal?	13
1.3.7 What is a spectrum, what is displayed on the x and y axis?	13
1.3.8 What is the role of the spectrum’s phase values of a signal?	14
1.3.9 When calculating a DFT, what parameters influence the frequency resolution of the spectrum?	14
1.3.10 Give examples for frequency-domain features of an audio signal!	14
1.4 Time-Frequency Domain Features	15
1.4.1 What are the problems when calculating the Fourier transform on a longer piece of music?	15
1.4.2 How can we maintain time information when performing frequency analysis of a signal?	15
1.4.3 Explain in your own words what the STFT is and how it is calculated!	15
1.4.4 What is the advantage of using the STFT instead of a simple Fourier Transform? .	16
1.4.5 What is the role of the window function in the context of the STFT?	16
1.4.6 What is a spectrogram and what values are shown on the x, y, and z axis? . .	17

1.4.7	Sketch a spectrogram of a signal that contains multiple short sine waves with different frequencies!	17
1.4.8	Give examples for time-frequency domain features!	17
1.5	Psychoacoustics	18
1.5.1	Explain why psychoacoustics are relevant for audio signal processing!	18
1.5.2	How are psychoacoustic scales created?	18
1.5.3	Give examples for psychoacoustic pitch scales! What are their characteristics?	18
1.5.4	How do psychoacoustics influence loudness perception?	19
1.5.5	How can we apply psychoacoustic loudness scaling to spectrograms?	19
1.5.6	How can we apply psychoacoustic frequency scales to spectrograms?	19
2	Onset Detection	20
2.1	Notes and Onsets	20
2.1.1	What are the (small-scale) structures that we can find in music?	20
2.1.2	How can we call the individual musical events in music?	20
2.1.3	Phases of the ADSR Model and Their Relation to Playing a Note on an Instrument	20
2.1.4	List basic properties of musical notes that we discussed in the lecture!	21
2.1.5	What is the relation of onsets and transients of musical notes?	21
2.1.6	Explain the difference between percussive, harmonic, and percussive-harmonic sounds!	21
2.1.7	Discuss identifying onsets for percussive, percussive-harmonic, and purely harmonic sounds and the challenges for each sound type!	22
2.2	Onset Detection - I. Signal Pre-Processing	22
2.2.1	What is the goal of onset detection?	22
2.2.2	What are challenges of onset detection in real-world signals?	23
2.2.3	Why is onset detection an important task / what can onsets be used for?	23
2.2.4	What are the steps in a typical processing pipeline for onset detection?	23
2.2.5	What is the goal of the signal preprocessing step for onset detection?	24
2.2.6	Give examples for signal preprocessing for onset detection!	24
2.3	Onset Detection - II. Detection Functions	25
2.3.1	What is an onset detection function?	25
2.3.2	What are the requirements for an onset detection function?	25
2.3.3	Give examples for onset detection functions!	25
2.3.4	Explain the fundamental idea of the phase-deviation onset detection function in your own words!	25
2.3.5	How is the spectral flux onset detection function calculated? No equation, only steps of algorithm.	26
2.3.6	Name three different types of audio signal contents and explain the challenges for onset detection for them!	26
2.4	Onset Detection - III. Peak Picking	27
2.4.1	What is the goal of peak picking in the context of onset detection?	27
2.4.2	What are the challenges of peak picking in the context of onset detection?	27
2.4.3	What are the processing steps for the peak picking method we discussed in the lecture?	27
2.5	Performance Evaluation	28
2.5.1	What kind of annotations do we need to evaluate an onset detection method?	28

2.5.2	What is a typical tolerance window for onset detection evaluation?	28
2.5.3	What evaluation metrics can be used for onset detection?	28
2.5.4	What are true positives, false positives, and false negatives in the context of onset detection evaluation?	28
2.5.5	What are true negatives, and why are they not used for onset detection evaluation?	28
2.5.6	Explain precision, recall, and f-measure. How are they calculated?	29
2.6	State of the Art Examples	29
2.6.1	Give an example of a concrete implementation for an onset detector!	29
2.6.2	Give an example for a vibrato compensation method for onset detection!	30
2.6.3	What is the difference between offline, online, and real-time processing?	30
2.6.4	What difficulties bring online processing in the context of onset detection?	30
2.7	Onset Detection via ML	30
2.7.1	Why is machine learning a suitable technique for onset detection?	30
2.7.2	Describe a typical ML pipeline for onset detection!	30
2.7.3	How can we represent audio input for CNNs?	31
2.7.4	What are the advantages of using ML for onset detection?	31
2.7.5	What are the drawbacks of using ML for onset detection?	31
3	Beat Tracking	32
3.1	Introduction	32
3.1.1	What are beats, downbeats, and measures, in the context of music?	32
3.1.2	What is tempo (in the context of music) and how can we quantify it?	32
3.1.3	Explain the difference between local and global tempo of a music track.	32
3.1.4	What is the goal of beat tracking?	32
3.2	Tempo Estimation	33
3.2.1	What is the goal of tempo estimation?	33
3.2.2	What are inter-onset-interval histograms, and how can they be used for tempo estimation?	33
3.2.3	What is autocorrelation, and how can it be used for tempo estimation?	33
3.2.4	What are octave errors in the context of tempo estimation?	34
3.2.5	What is a tempogram?	34
3.3	Beat Tracking	34
3.3.1	What is beat phase detection, in the context of beat tracking?	34
3.3.2	Describe a method to detect the beat phase for beat picking.	34
3.3.3	Explain the greedy beat tracking algorithm discussed in the lecture.	35
3.3.4	Explain the multi-agent tracking approach discussed in the lecture.	35
3.3.5	Explain the machine learning setup for beat tracking.	36
3.3.6	Why are RNNs a suitable architectural choice for beat tracking?	36
3.3.7	What is the task of the RNN in the RNN beat tracker discussed in the lecture?	36
3.3.8	What is the task of the DBN in the RNN beat tracker discussed in the lecture?	36
3.4	Evaluation	36
3.4.1	Which metrics can we use for beat tracking evaluation?	36
3.4.2	In what range are tolerances for beat positions typically?	37
3.4.3	Why are tolerances for beat tracking higher than for onset detection?	37
4	F0 and Notes	38

4.1	Pitch and F0	38
4.1.1	Properties of Musical Notes Mentioned in the Lecture	38
4.1.2	Pitch and Its Relationship to the Fundamental Frequency of a Note	38
4.1.3	Harmonics or Harmonic Partials of a Note	38
4.1.4	Fundamental Frequency of a Note	38
4.1.5	Difference Between Percussive, Harmonic, and Harmonic-Percussive Notes	38
4.1.6	Visualization Method for Monophonic F0 Estimation/Annotations	38
4.1.7	Visualization Methods for Polyphonic F0 Estimation/Annotations	39
4.1.8	Difference Between Single- and Multi-F0 Estimation	39
4.1.9	Three Levels for F0 Estimation Discussed in the Lecture	39
4.2	Evaluation	40
4.2.1	Evaluation for F0 Estimation on Frame-Level	40
4.2.2	Evaluation for F0 Estimation on Note-Level	40
4.2.3	Evaluation for F0 Estimation on Stream-Level	40
4.2.4	Issues with Frame-Level F0 Estimation	40
4.2.5	Metrics for Frame-Level F0 Estimation	41
4.2.6	Issues with Note-Level F0 Estimation	41
4.2.7	Properties Evaluated for Note-Level F0 Estimation	41
4.3	Fundamental Frequency Estimation	41
4.3.1	Which simple/naive methods could be used for F0 estimation?	41
4.3.2	Explain the typical problems of simple F0 estimation methods!	41
4.3.3	Explain the processing steps for the YIN F_0 estimation method!	42
4.3.4	What is the difference between YIN and pYIN?	43
4.4	F0 Estimation in the Frequency Domain	44
4.4.1	What is the fundamental idea behind F_0 estimation in the frequency domain?	44
4.4.2	List some of the problems of F_0 estimation in the frequency domain!	44
4.4.3	How does the SHAPE F_0 estimation method work, in principle?	45
4.5	Multiple F0 Estimation	46
4.5.1	Explain the difference between single and multiple- F_0 estimation!	46
4.5.2	What are the challenges of multiple- F_0 estimation?	46
4.5.3	Explain the three categories for multiple- F_0 estimation we discussed in the lecture!	47
4.6	Neural Network Based Methods	48
4.6.1	Explain how a CNN can be used to perform multi-f0 estimation!	48
4.6.2	What issues can arise when doing frame-wise multiple-f0 estimation?	48
4.6.3	What is the task of the HMM post-processing for the CNN multiple-f0 estimation approach, we discussed in the lecture?	48
4.6.4	Why does explicit modeling of note phases (ADSR) improve performance for multiple-f0 estimation?	49
5	Music Classification and Semantic Annotation	50
5.1	Semantic Description	50
5.2	Musical Genre	50
5.3	Why Genre Classification?	50
5.4	Genre Classification	51
5.5	Case Study: k-NN for Genre Classification	51
5.6	Why NOT Genre Classification?	51

5.7	Tags	52
5.8	Games With A Purpose (GWAP)	52
5.8.1	ESP Game	52
5.8.2	TagATune	52
5.8.3	MagnaTagATune Dataset	53
5.9	Audio Feature Extraction	53
5.9.1	Mel Frequency Cepstral Coefficients (MFCCs)	53
5.9.2	The Mel Scale	53
5.9.3	Bag-of-Frames Modeling	53
5.10	Deep Learning Approaches	54
5.11	Deep Learning for Music Tagging and Classification	54
5.11.1	End-to-End Learning for Tags	54
5.11.2	CNN Architectures and Design	54
5.11.3	Comparison of Input Representations and CNN Models	55
5.11.4	Results from Choi et al. (2016)	55
5.11.5	Experimental Setup and Analysis	55
6	Multimodality	56
6.1	Diverse Music Data Sources	56
6.2	Collaborative Filtering (CF)	56
6.3	Factors Hidden in the Data	56
6.4	Audio Content Analysis	57
6.5	Recap: Audio Content Analysis (2017 State-of-the-Art)	57
6.6	Text Analysis Methods (Basic IR)	58
6.7	Textual Sources of Music Similarity	58
6.8	Multimodal Approaches	58
6.9	Feedback-Transformed Content	59
6.10	Beyond Items	59
6.11	Listener Background	59
6.12	Listener Context	59
6.13	Services, Data Bias and Limitations	60
7	Drum Transcription	61
7.1	Introduction	61
7.1.1	Explain the automatic drum transcription (ADT) task!	61
7.1.2	What are the differences between drum transcription and harmonic instrument transcription?	61
7.1.3	What kind of data is needed for automatic drum transcription (training and evaluation)?	62
7.1.4	Which metrics are used for ADT evaluation?	62
7.2	Methods Overview	63
7.2.1	What are the three categories of ADT methods we discussed in the lecture? .	63
7.2.2	Explain the segment and classify approach for ADT!	63
7.2.3	Explain the separate and detect approach for ADT!	63
7.2.4	Explain the activation function based approach for ADT!	63
7.3	Non-Negative Matrix Factorization	64
7.3.1	Explain in your own words the algorithm for non-negative matrix factorization!	64

7.3.2	What are the challenges that come with NMF for automatic drum transcription?	65
7.3.3	Name variants for NMF in the context of ADT, what are the differences?	65
7.4	NN Based Methods	67
7.4.1	Explain the setup for neural network based ADT!	67
7.4.2	Explain the difference between the multi-task and single-task NN approach for ADT!	67
7.4.3	What are questions to answer when trying to choose a model for a given task?	67
7.4.4	Explain the differences between dense, recurrent, convolutional, and dilated-conv neural networks! What are the advantages and disadvantages of each architecture category?	68
7.4.5	Explain how a CNN/RNN/TCN can be used to perform automatic drum transcription! How is the spectrogram processed in the context for each model?	68
8	Audio Representation Learning	69
8.1	Introduction to Audio Representation Learning	69
8.2	Machine Learning Paradigms	69
8.2.1	Rule-Based Systems	69
8.2.2	Classic Machine Learning	69
8.2.3	Representation Learning	70
8.2.4	Deep Learning	70
8.3	Audio Feature Extraction	70
8.4	Signal Processing for Audio Representations	70
8.5	Approaches to Audio Representation Learning	71
8.5.1	Linear Methods	71
8.5.2	Nonlinear Methods	71
8.5.3	Deep Representation Learning	71
8.6	Deep Learning Methods for Audio	71
8.6.1	Autoencoders (AEs)	71
8.6.2	Supervised Learning for Representation Extraction	71
8.6.3	Triplet Networks	71
8.7	Challenges in Audio Representation Learning	72
8.8	Applications of Audio Representation Learning	72
9	Acoustic Scene Classification	73
9.1	Introduction to Acoustic Scene Classification (ASC)	73
9.2	Scene Classification vs. Sound Event Detection	73
9.3	Applications of ASC and SED	73
9.3.1	Health and Medical Sector	73
9.3.2	Security and Surveillance	73
9.3.3	Ecosystem Monitoring	73
9.3.4	Other Applications	73
9.4	Challenges in Acoustic Scene Classification and Sound Event Detection	74
9.4.1	Generalization Challenges	74
9.4.2	Data Augmentation to Address Challenges	74
9.5	Data for ASC and SED	74
9.5.1	Ground Truth Label Types	74
9.5.2	Data Annotation Techniques	74

9.6	Approaches to ASC and SED	74
9.6.1	Machine Learning for Acoustic Scene Classification	74
9.6.2	Deep Learning for ASC and SED	75
9.7	Evaluation Campaigns and Datasets	75
9.7.1	Key Evaluation Campaigns	75
9.7.2	Popular Datasets	75

1 Fundamentals

1.1 Basics

1.1.1 What are the differences between sound, audio, and music?

- **Sound:** Generated by a **vibrating object**, causing **displacements** and **oscillations** of air molecules. Measured in **Hertz (Hz)**, which represents cycles per second.
- **Audio:** The transmission, reception, or reproduction of **sound** within the human hearing range (20 Hz to 20 kHz). Includes dynamic and tonal information essential for accurate sound reproduction.
- **Music:** Structured arrangement of **audio events** in terms of **time** and **pitch**, often with artistic or emotional intent.

1.1.2 What is the frequency range for audio?

- The frequency range for human hearing is typically 20 Hz to 20 kHz.

1.1.3 Give examples for objective properties in music!

- **Pitch:** Frequency of the sound wave.
- **Duration:** Time the sound is held.
- **Loudness:** Perceived intensity or amplitude of the sound wave.
- **Timbre:** Unique quality or color of the sound.

1.1.4 What is an audio waveform and how can it be created from sound?

- An **audio waveform** is a representation of sound, with **time** on the x-axis and **amplitude** on the y-axis.
- It is created by converting sound pressure variations into an electrical signal, typically using a microphone.

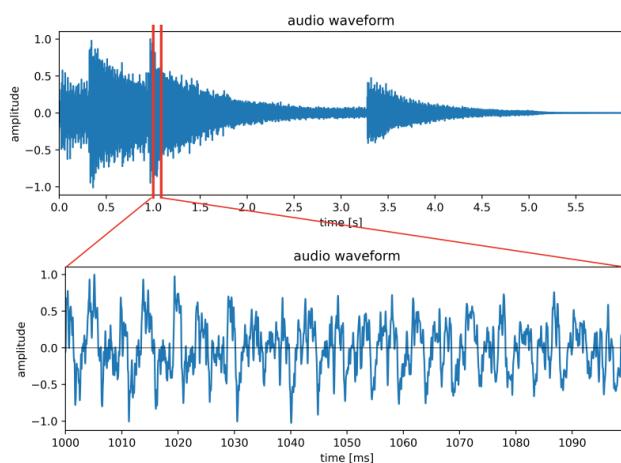


Figure 1: Audio Waveform

1.1.5 What is ADC and what steps does it consist of?

- **Analog-to-Digital Conversion (ADC)** converts analog signals into digital data.
- Steps include:
 1. **Sampling:** Capturing the signal at discrete time intervals.
 2. **Quantization:** Mapping amplitude values to discrete levels.

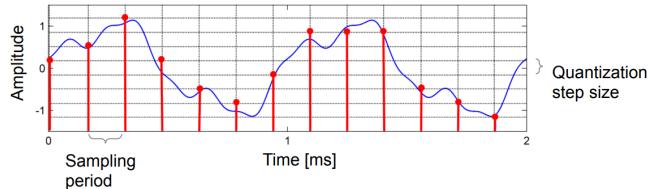


Figure 2: ADC example

3. **Encoding:** Representing the quantized values as digital data.

1.1.6 What influence do sampling frequency and bit depth for ADC have?

- **Sampling Frequency:** Determines the temporal resolution and the highest frequency that can be accurately captured (**Nyquist frequency**).
- **Bit Depth:** Affects the amplitude resolution and dynamic range. Higher bit depth reduces quantization error.

1.1.7 What are the requirements for the sampling frequency for ADC, when sampling a signal?

- The sampling frequency must be at least twice the highest frequency in the signal, as per the **Nyquist–Shannon Sampling Theorem**. Frequencies above half of sampling frequency (Nyquist frequency) cannot be captured.
- Common audio sampling rates: 44.1 kHz (CD quality), 48 kHz, 96 kHz.

1.1.8 What does PCM stand for?

- PCM stands for **Pulse-Code Modulation**.

1.1.9 Explain how PCM encodes an audio signal!

- PCM encodes an audio signal by:
 1. Sampling the analog signal at equidistant intervals.
 2. Quantizing the amplitude of each sample to a finite set of values.
 3. Encoding the quantized values as binary data.
- PCM is the foundation of most digital audio formats.

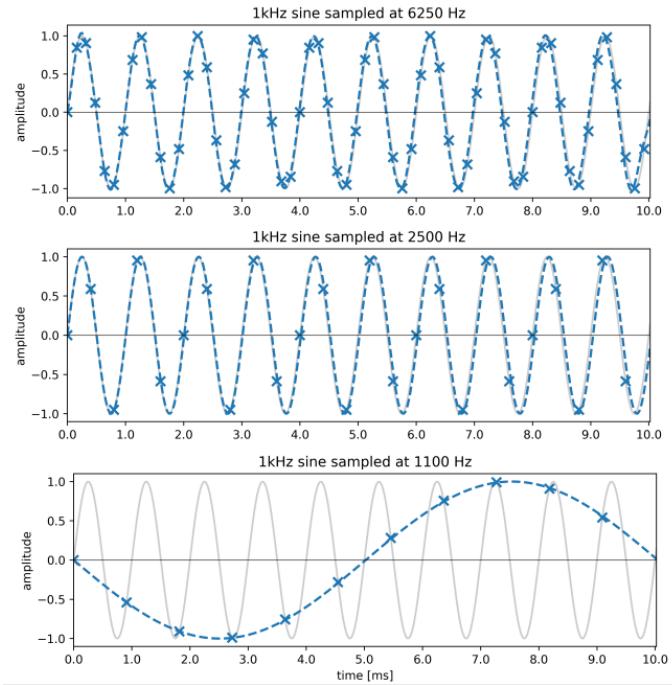


Figure 3: Nyquist–Shannon Sampling Theorem

1.2 Audio Features and Signal Framing

1.2.1 What are (audio-)features in the context of signal processing?

- **Audio features** are numerical representations or summaries of the properties of an audio signal.
- They are used to extract meaningful information from audio signals for analysis or classification.
- Features can be categorized as:
 - **Time-domain features**: Derived directly from the audio waveform.
 - **Frequency-domain features**: Derived from the signal's spectral representation (e.g., Fourier Transform).
 - **Time-frequency domain features**: Combine time and frequency information (e.g., spectrogram).

1.2.2 What are examples for basic time-domain audio features?

- **Root Mean Square (RMS) Energy**: Measures the average energy of the signal; correlates with perceived loudness.
- **Zero Crossing Rate (ZCR)**: Counts the number of times the signal crosses the zero-amplitude line; indicates signal noisiness or pitch.
- **Amplitude Envelope**: Tracks the maximum amplitude within short time frames; useful for analyzing dynamics.

1.2.3 What are problems of global audio features, especially in the context of music?

- **Limited Temporal Information:** Global features represent the entire signal with a single value, ignoring temporal changes.
- **Dynamic Nature of Music:** Music often changes in pitch, rhythm, and timbre over time; global features cannot capture these variations.
- **Dependency on Signal Length:** Some global features may vary with the length of the signal, reducing their reliability for comparison.

1.2.4 Explain in your own words what signal framing is, including the relevant parameters!

- **Signal framing** is the process of dividing an audio signal into short, overlapping segments called **frames**.
- Key parameters:
 - **Frame length:** Number of samples in each frame (e.g., 512–8192 samples, corresponding to 11–186 ms at 44.1 kHz).
 - **Hop size:** Number of samples between the start of consecutive frames; determines the overlap.
 - **Overlap:** Difference between frame length and hop size; ensures smooth transitions between frames.

The main idea is: `overlap = frame_len - hop_size`

- Purpose: To capture short-time properties of the signal while preserving time-local information.

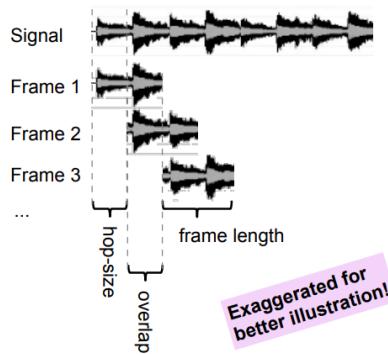


Figure 4: Signal Framing

1.2.5 How can signal framing be used to calculate time-local audio features?

- The signal is divided into frames, and features (e.g., ZCR, RMS, spectral features) are calculated for each frame individually.
- By analyzing each frame, variations in features over time can be captured, revealing the dynamic nature of the signal.

1.2.6 What is the advantage of time-local audio features?

- Time-local features preserve information about changes in the signal over time, crucial for dynamic audio content like music.
- They enable tasks such as beat tracking, onset detection, and speech analysis, where temporal variations are critical.

1.2.7 Give examples of time-local audio features and explain what information they reveal!

- **Zero Crossing Rate (ZCR):**
 - Measures the rate at which the signal crosses the zero line.
 - Reveals information about pitch or noisiness of the signal.

- **Short-Time Energy (RMS):**
 - Measures the signal's power within a frame.
 - Indicates loudness and energy dynamics.

- **Spectral Centroid:**
 - Represents the center of gravity of the spectrum.
 - Indicates whether the sound is dominated by high or low frequencies.

1.2.8 Explain the process of calculating a time-local feature like ZCR!

- **Zero Crossing Rate (ZCR):**
 1. Divide the signal into overlapping frames.
 2. For each frame, calculate the number of times the signal changes sign (crosses zero).
 3. Normalize the count by the frame length to obtain the ZCR for that frame:

$$ZCR_t = \frac{1}{2} \sum_{k=tK}^{(t+1)K-1} |\operatorname{sgn}(s(k)) - \operatorname{sgn}(s(k+1))|$$

where $s(k)$ is the amplitude of the k -th sample, and K is the frame length.

- **Interpretation:** A high ZCR indicates a noisy or high-frequency signal, while a low ZCR suggests smoother or lower-frequency content.

1.3 Frequency Domain Features

1.3.1 How does audio signal frequency analysis relate to the human hearing process?

The human auditory system analyzes audio signals in terms of frequency, much like how frequency-domain analysis works in signal processing. Within the cochlea, specific regions respond to distinct frequency ranges, decomposing complex sounds into simpler frequency components. Similarly, frequency analysis methods such as the Fourier Transform mimic this process by breaking down audio signals into their frequency components, helping to understand characteristics like pitch, timbre, and harmonics.

1.3.2 Explain, in your own words, what the Fourier transform is!

The Fourier Transform is a mathematical method that converts a signal from its **time-domain representation** (amplitude over time) into its **frequency-domain representation** (amplitude and phase over frequency). It essentially breaks a signal down into a sum of sinusoidal waves of varying frequencies, amplitudes, and phases. This makes it possible to identify the underlying components of complex signals, which is invaluable for tasks like pitch detection and spectral analysis.

1.3.3 Sketch a sine-frequency with 10Hz and the magnitude spectrum of this signal!

1.3.4 Sketch a signal that is a mixture of two sines and its magnitude spectrum!

1.3.5 Explain amplitude, frequency, and phase for a sinusoid!

A sinusoid is fully described by its amplitude, frequency, and phase.

- The amplitude represents the maximum strength or height of the wave, correlating with how loud the sound appears.
- The frequency defines how often the sinusoid oscillates per second, measured in Hertz (Hz), and directly relates to the pitch of the sound.
- The phase describes the relative starting point of the wave, measured in radians or degrees, and influences how the sinusoid aligns with others when combined.

1.3.6 What is a “frequency-domain representation” of an audio signal?

A frequency-domain representation describes an audio signal in terms of its frequency components rather than its time-dependent waveform. It reveals how much energy is present at each frequency, making it easier to analyze harmonic content, detect dominant pitches, and study spectral characteristics of the signal.

1.3.7 What is a spectrum, what is displayed on the x and y axis?

A spectrum is a graphical representation of a signal's frequency content. The x-axis typically shows the frequencies in Hertz, while the y-axis indicates the magnitude, representing the energy or amplitude of the signal at each frequency. It is a powerful tool for visualizing and analyzing the frequency components of an audio signal.

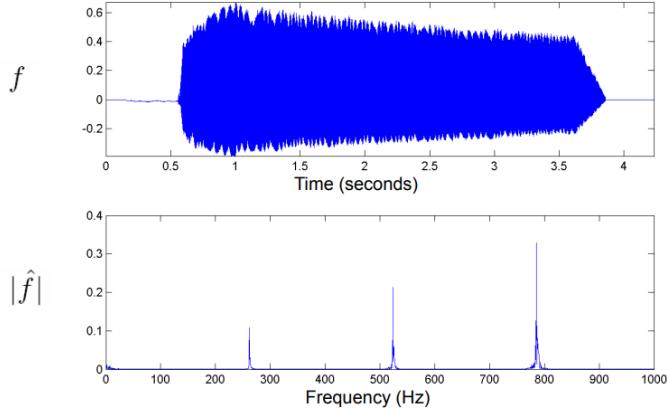


Figure 5: Time-domain to Frequency-domain example

1.3.8 What is the role of the spectrum's phase values of a signal?

The phase values in a spectrum indicate the relative timing of frequency components. While the magnitude spectrum shows the strength of each frequency, the phase spectrum is essential for reconstructing the original signal accurately. Phase affects how the frequency components combine in time, influencing the overall shape and alignment of the waveform.

1.3.9 When calculating a DFT, what parameters influence the frequency resolution of the spectrum?

The frequency resolution of a Discrete Fourier Transform (DFT) depends on two key parameters:

1. The sampling frequency
2. The number of samples

A higher number of samples results in smaller frequency intervals, improving resolution. Similarly, the sampling frequency determines the range of frequencies that can be analyzed. Together, these parameters control the ability to distinguish between closely spaced frequencies in the spectrum.

1.3.10 Give examples for frequency-domain features of an audio signal!

Frequency-domain features include:

1. **Spectral centroid**, which represents the center of gravity of the spectrum and indicates whether a sound is dominated by high or low frequencies
2. **Bandwidth**, which describes the range of significant frequencies in the signal.

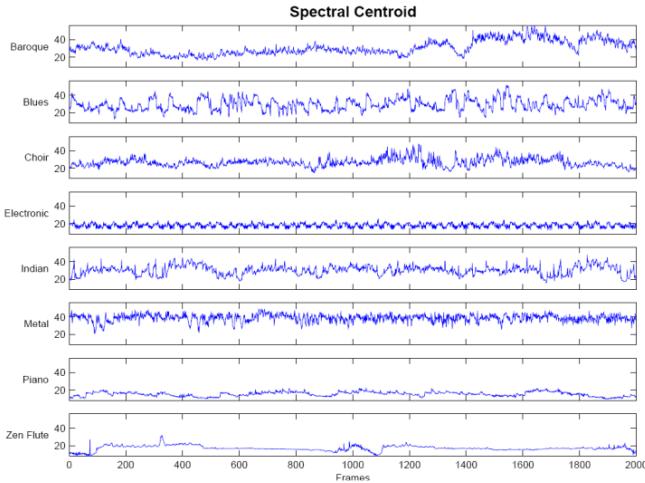


Figure 6: Spectral Centroid example

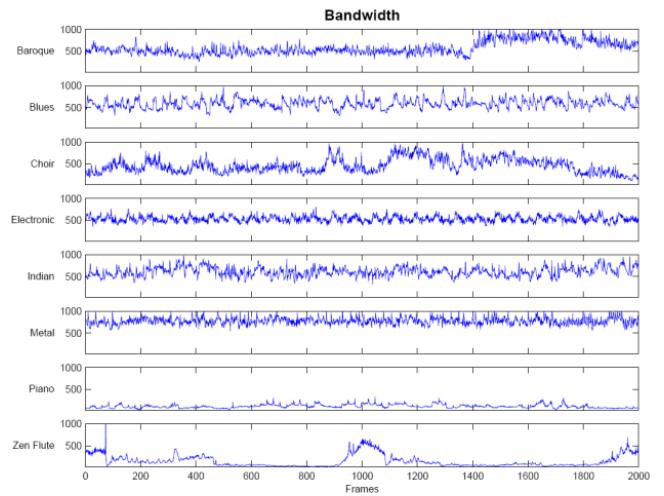


Figure 7: Bandwidth Example

1.4 Time-Frequency Domain Features

1.4.1 What are the problems when calculating the Fourier transform on a longer piece of music?

The Fourier Transform provides information about the frequency components of a signal but averages this information over the entire duration of the signal, losing time-local information.

Music and other audio signals are dynamic, with changing frequencies, harmonics, and rhythms over time, which the Fourier Transform cannot capture.

For long signals, the assumption of **stationarity** (unchanging frequency components) is invalid, making the analysis less meaningful.

1.4.2 How can we maintain time information when performing frequency analysis of a signal?

By dividing the signal into short, overlapping segments (frames) and performing frequency analysis on each segment individually. This approach, called **the Short-Time Fourier Transform (STFT)**, allows frequency components to be analyzed in small time intervals, maintaining temporal resolution.

1.4.3 Explain in your own words what the STFT is and how it is calculated!

The Short-Time Fourier Transform (STFT) is a technique that combines time and frequency analysis by dividing a signal into short overlapping frames and performing a Fourier Transform on each frame.

The steps are:

1. Divide the signal into overlapping frames of fixed length.
2. Multiply each frame by a window function to reduce edge effects.
3. Perform the Fourier Transform on each windowed frame to obtain its frequency content.

The result is a two-dimensional representation (time vs. frequency) of the signal.

1.4.4 What is the advantage of using the STFT instead of a simple Fourier Transform?

The STFT provides both time and frequency information, allowing analysis of non-stationary signals. It enables tracking of changes in frequency content over time, making it particularly useful for analyzing music, speech, and other dynamic signals.

1.4.5 What is the role of the window function in the context of the STFT?

The window function determines how much of the signal is included in each frame and reduces artifacts caused by abrupt frame boundaries.

Common window functions, like the Hann or Hamming window, ensure smooth transitions between frames, reducing spectral leakage.

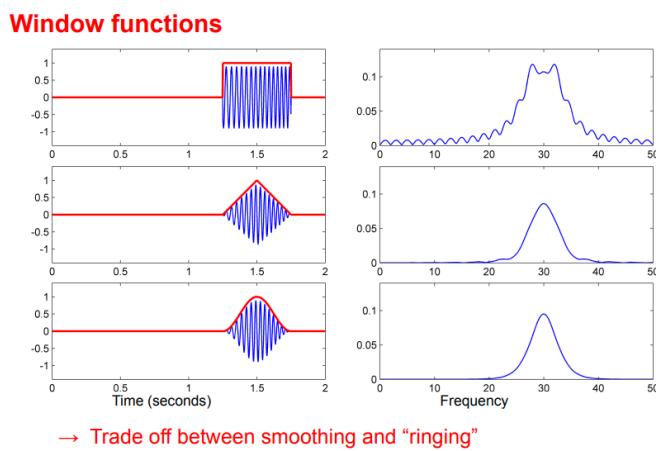


Figure 8: STFT Window Functions

1.4.6 What is a spectrogram and what values are shown on the x, y, and z axis?

A spectrogram is a visual representation of a signal's time-frequency content.

1. **x-axis:** Time (seconds).
2. **y-axis:** Frequency (Hertz).
3. **z-axis (color intensity):** Magnitude or energy of the signal at each frequency and time.

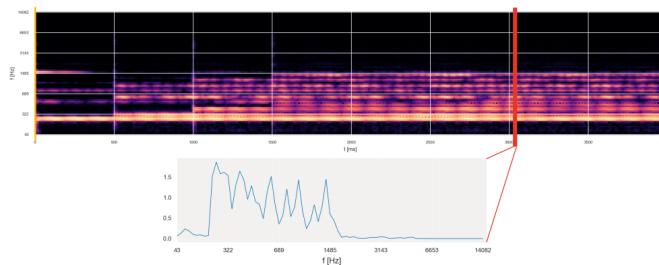


Figure 9: Spectrogram Example

1.4.7 Sketch a spectrogram of a signal that contains multiple short sine waves with different frequencies!

1.4.8 Give examples for time-frequency domain features!

- **Spectral Centroid:** Indicates the "center of gravity" of the spectrum at each time frame.
Description: The center of gravity of the magnitude spectrum of the DFT, i.e., the frequency (band) region where most of the energy is concentrated.
Remarks:
 - Used as a measure of sound sharpness (strength of high-frequency energy).
 - Sensitive to low-pass filtering (downsampling) as the high-frequency bands are given more weight.
 - Sensitive to white noise (for the same reason).
- **Spectral Bandwidth:** Represents the spread of the frequency spectrum around the centroid.
Remarks:
 - The average bandwidth of a piece of music may serve as an indicator of aggressiveness.
 - Does not provide information about the perceived rhythmic structure.
 - Not suitable to distinguish different parts of a piece of music (e.g., vocal part in a metal piece may not be visible).
- **Spectral Flux:** Measures changes in the spectrum over time.
- **Mel-Frequency Cepstral Coefficients (MFCCs):** Capture perceptually meaningful features derived from the spectrogram.

1.5 Psychoacoustics

1.5.1 Explain why psychoacoustics are relevant for audio signal processing!

Psychoacoustics studies how humans perceive sound, helping to design systems that align with human hearing capabilities.

By understanding perception, we can prioritize relevant audio information (e.g., important frequencies or loudness levels) and optimize compression, synthesis, and playback systems.

1.5.2 How are psychoacoustic scales created?

Psychoacoustic scales are derived from experiments that map physical sound properties (e.g., frequency, intensity) to perceived characteristics (e.g., pitch, loudness).

Examples include:

- Mapping frequency to perceptual pitch (e.g., Mel scale).
- Relating sound pressure to perceived loudness (e.g., Equal Loudness Contours).

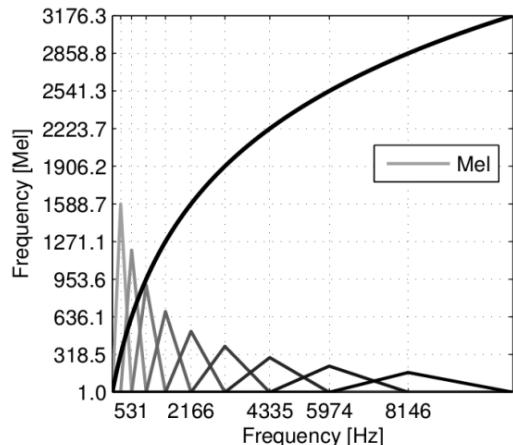


Figure 10: Mel Scale Example

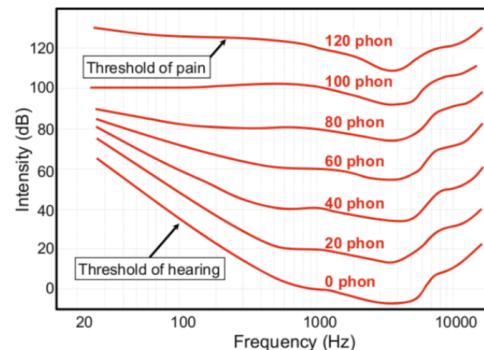


Figure 11: Equal Loudness Contour

1.5.3 Give examples for psychoacoustic pitch scales! What are their characteristics?

Psychoacoustic pitch scales represent how humans perceive pitch, focusing on the nonlinear nature of auditory perception. Key scales include:

- **Mel Scale:** Reflects perceived pitch with equal spacing corresponding to equal perceptual differences. It is linear at low frequencies and logarithmic at high frequencies. Typically, 40 Mel-frequency bins are used, spaced equally on this scale.
- **Bark Scale:** Divides the frequency spectrum into critical bands, approximating human auditory filters. It is useful for studying auditory masking and sound interactions.
- **Cent Scale:** A logarithmic pitch scale where 1200 cents equal one octave. It provides precise representation of small pitch intervals.

- **Semitone Scale:** Maps frequencies to semitone intervals with a reference frequency (usually 440 Hz). Uses a filter bank with triangular filters centered on semitone frequencies, each spanning 100 cents.

- **dB Loudness Scaling:** A logarithmic scale for perceived loudness. Sound pressure levels in decibels are calculated as:

$$L_p = 20 \cdot \log_{10} \left(\frac{p}{p_0} \right)$$

For spectrograms, power values ($|S|^2$) are expressed as:

$$S_{dB} = 10 \cdot \log_{10} |S|^2$$

This allows relative loudness visualization when the reference value is omitted.

1.5.4 How do psychoacoustics influence loudness perception?

Loudness perception depends on frequency, with humans being more sensitive to mid-frequencies (1–5 kHz) and less sensitive to very low or very high frequencies.

The **Equal Loudness Contours** describe this non-linear perception, showing how different sound pressures are required for equal perceived loudness at various frequencies.

1.5.5 How can we apply psychoacoustic loudness scaling to spectrograms?

- Use logarithmic scaling (e.g., decibels) to reflect human perception of loudness.
- Normalize intensity values to emphasize perceptually important parts of the spectrum.

1.5.6 How can we apply psychoacoustic frequency scales to spectrograms?

- Transform the linear frequency axis into a psychoacoustic scale such as Mel or Bark.
- Apply a filter bank (e.g., Mel filter bank) to group and smooth frequency bins according to perceptual relevance.
- The resulting spectrogram reflects how humans perceive the frequency content of the signal.



Figure 12: Filterbank example

2 Onset Detection

2.1 Notes and Onsets

2.1.1 What are the (small-scale) structures that we can find in music?

Small-scale structures in music include:

- Events, individual notes.
- Chords
- Rhythmic patterns

These elements create the detailed framework of a piece, contributing to its texture and expressiveness.

2.1.2 How can we call the individual musical events in music?

Individual musical events are referred to as:

- *Notes* when describing pitch and duration
- *Onsets* when focusing on the timing of their initiation
- *Attacks* when emphasizing the initial phase of sound production.

2.1.3 Phases of the ADSR Model and Their Relation to Playing a Note on an Instrument

The ADSR model describes the envelope of a sound in four distinct phases, which correspond to how the amplitude of a sound changes over time. These phases are crucial for understanding how musical instruments produce and shape their sound.

- **Attack (transient):** The initial phase where the sound's amplitude increases rapidly. For example, on a piano, this corresponds to the moment the hammer strikes the string, creating a sudden onset of vibration and sound.
- **Decay:** Following the attack, this phase sees the amplitude decrease to a steady level. In a piano, this corresponds to the string's vibration stabilizing after the initial strike.
- **Sustain (resonance):** This is the phase where the sound is held at a relatively constant amplitude. The duration of this phase depends on the instrument and playing technique, such as bowing a violin string continuously or holding down a piano key.
- **Release:** The final phase where the sound fades to silence. On a piano, this happens when the key is released, and the damper stops the string's vibration.

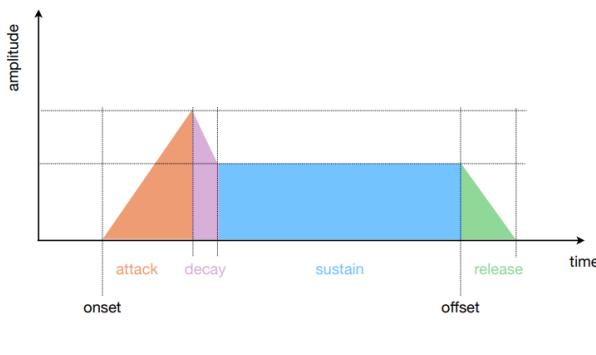


Figure 13: ADSR Model

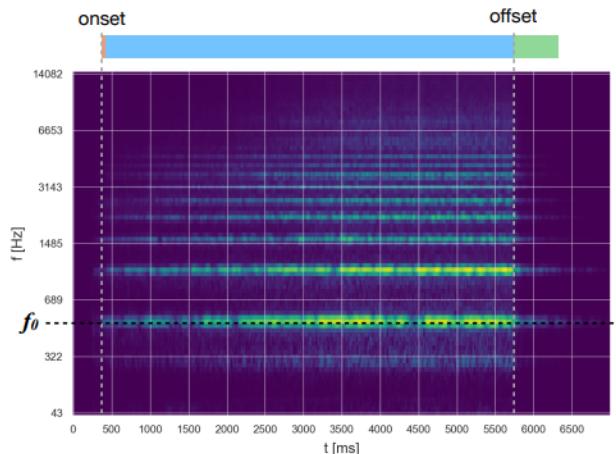


Figure 14: ADSR Spectrogram Example

2.1.4 List basic properties of musical notes that we discussed in the lecture!

Musical notes are characterized by:

- Pitch (frequency)
- Duration
- Intensity (loudness)
- Timbre
- Onset time

These properties determine how notes are perceived and how they interact in a musical context.

2.1.5 What is the relation of onsets and transients of musical notes?

The **onset** marks the **beginning of a note**, corresponding to the point in time when sound energy significantly increases. The **transient** refers to the brief, non-stationary period at the start of the note, often encompassing the attack phase, where the spectral content changes rapidly.

2.1.6 Explain the difference between percussive, harmonic, and percussive-harmonic sounds!

- **Percussive sounds** are transient-rich and characterized by sharp, short attacks, such as drums or clapping.
- **Harmonic sounds** are sustained and exhibit stable periodic frequencies, like a bowed violin or a singing voice.
- **Percussive-harmonic sounds** combine both qualities, featuring distinct transients followed by sustained harmonic content, such as a piano or a plucked guitar string.

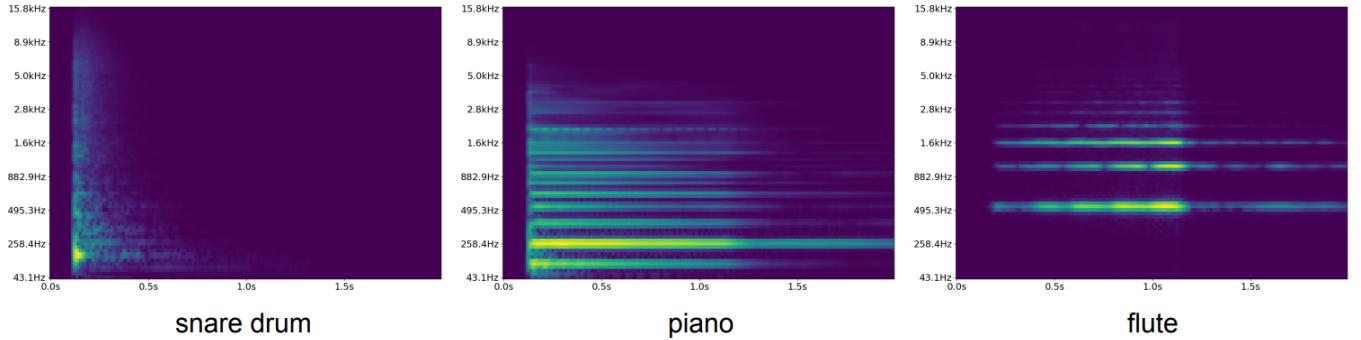


Figure 15: Different onsets example

2.1.7 Discuss identifying onsets for percussive, percussive-harmonic, and purely harmonic sounds and the challenges for each sound type!

Identifying onsets in **percussive sounds** is relatively straightforward due to their clear and sudden amplitude changes. For **harmonic sounds**, the challenge lies in detecting gradual energy increases or subtle spectral shifts, as the attack phase is less defined. **Percussive-harmonic sounds** present intermediate difficulty, where transient detection is feasible, but distinguishing overlapping harmonic content may require additional spectral analysis techniques.

2.2 Onset Detection - I. Signal Pre-Processing

2.2.1 What is the goal of onset detection?

The goal of onset detection is to identify the precise moments in time when new musical events or sounds begin. These moments, called onsets, are critical for analyzing rhythmic structures, extracting features, and enabling tasks such as music transcription or beat tracking.

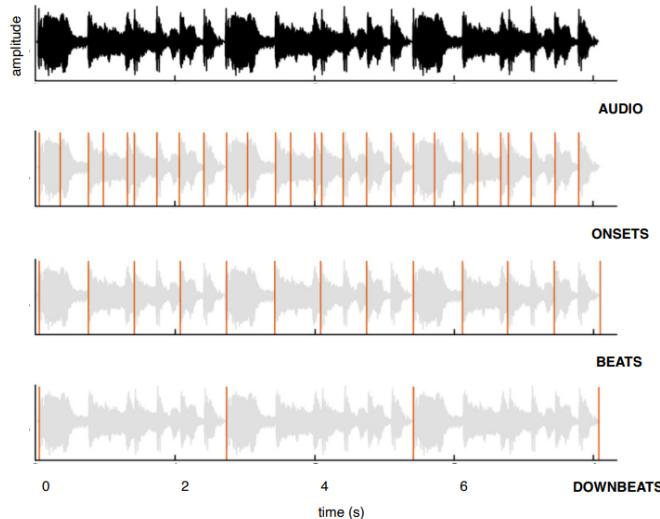


Figure 16: Onset, Beats and Downbeats detection

2.2.2 What are challenges of onset detection in real-world signals?

Challenges include:

1. **Variability in sound types**, such as percussive, harmonic, and mixed signals, each with unique characteristics.
2. **Overlapping sounds or polyphonic textures**, making it difficult to isolate individual onsets.
3. **Noise and distortions in real-world recordings**, which can obscure onset features.
4. **Gradual onsets**, particularly in harmonic or sustained sounds, where energy changes are subtle.

2.2.3 Why is onset detection an important task / what can onsets be used for?

Onset detection is fundamental for many audio processing tasks, such as:

1. **Automatic music transcription**.
2. **Beat and tempo analysis**.
3. **Audio segmentation** for music structure analysis.
4. **Feature extraction** for machine learning applications.
5. **Synchronization of audio and video** or audio effects in production.

2.2.4 What are the steps in a typical processing pipeline for onset detection?

1. Signal preprocessing to enhance relevant features and suppress noise.
2. Computation of an onset detection function to highlight energy or spectral changes.
3. Application of peak picking algorithms to locate the exact onset positions.

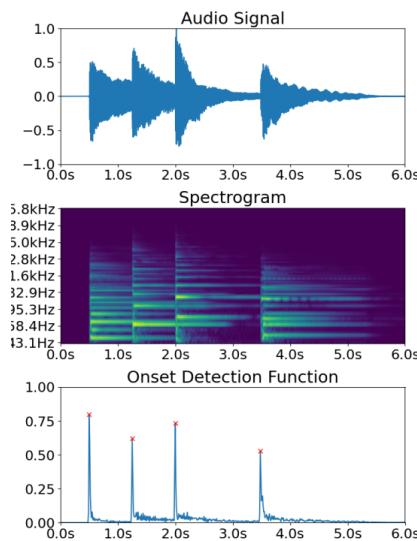


Figure 17: Pipeline

2.2.5 What is the goal of the signal preprocessing step for onset detection?

The goal of signal preprocessing is to prepare the audio signal for analysis by enhancing characteristics associated with onsets and reducing irrelevant or distracting components, such as noise or steady background signals.

2.2.6 Give examples for signal preprocessing for onset detection!

Examples of preprocessing techniques include:

- **High-pass filtering:** Removes low-frequency noise to focus on higher-frequency components relevant to onset detection.
- **Envelope extraction:** Captures the amplitude variations over time, which are often indicative of onsets in the signal.
- **Short-Time Fourier Transform (STFT):** Analyzes the time-frequency content, providing a spectrogram that reveals changes in energy over time.
- **Spectral smoothing:** Reduces noise and emphasizes abrupt changes in spectral content.
- **Adaptive whitening:** A technique to normalize the magnitudes $|X(n, k)|$ of each frequency bin based on past peak values. This is achieved using an iterative algorithm for causal (realtime) processing:

$$X_{\text{norm}}(n, k) = \frac{X(n, k)}{\max(r, m \cdot P(n - 1, k))}$$

Here:

- $X(n, k)$: Value of the STFT at frame index n and frequency bin index k .
- $P(n - 1, k)$: Past peak values.
- m : Memory coefficient controlling the influence of past values.
- r : Floor parameter to prevent division by very small values.

This approach dynamically adjusts to the signal's characteristics, enhancing its suitability for onset detection.

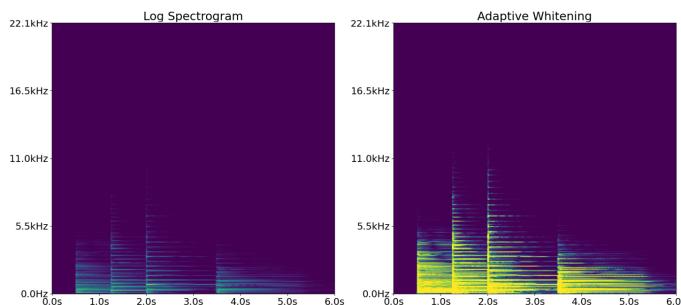


Figure 18: Adaptative Whitening example

2.3 Onset Detection - II. Detection Functions

2.3.1 What is an onset detection function?

An onset detection function is a time-domain signal derived from the audio input that emphasizes potential onsets by highlighting significant changes in energy, spectral content, or phase.

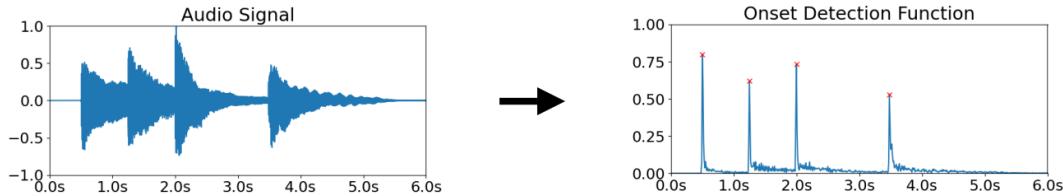


Figure 19: Example of onset detection function

2.3.2 What are the requirements for an onset detection function?

An effective onset detection function should:

1. Accurately capture changes in the audio signal that correspond to onsets.
2. Minimize false positives and negatives, even in noisy or complex signals.
3. Be computationally efficient for real-time applications.

2.3.3 Give examples for onset detection functions!

1. **Signal Envelope:** Tracks amplitude increases using a smoothed version of the signal. Improvements include using first-order differences or psychoacoustic "A-weighting." Simple but less reliable.
2. **High-Frequency Content:** Focuses on the broadband nature of transients by emphasizing high-frequency energy using weighted STFT magnitudes.
3. **Spectral Difference:** Measures changes between successive spectrogram frames using norms (e.g., L_1 or L_2). Spectral flux applies half-wave rectification for better results with logarithmic spectrograms.
4. **Phase Deviation:** Uses the second derivative of unwrapped phase from the STFT. Detects transients as deviations from constant phase increments, with "mean absolute phase deviation" being a common metric.

2.3.4 Explain the fundamental idea of the phase-deviation onset detection function in your own words!

The phase-deviation onset detection function identifies onsets by analyzing sudden and significant changes in the phase relationships of frequency components. These changes often occur when a new sound event disrupts the steady phase progression of ongoing frequencies.

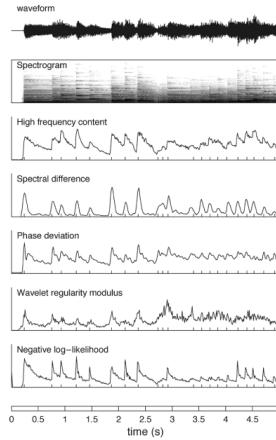


Figure 20: Example of onset detection functions

2.3.5 How is the spectral flux onset detection function calculated? No equation, only steps of algorithm.

1. Compute the Short-Time Fourier Transform (STFT) of the audio signal to obtain its spectral representation.
2. Calculate the magnitude spectrum for each time frame.
3. Measure the frame-to-frame changes in the magnitude spectrum by comparing consecutive frames.
4. Retain only positive changes (increases in energy) to focus on potential onsets.
5. Aggregate the changes across frequencies for each time frame to generate the detection function.

2.3.6 Name three different types of audio signal contents and explain the challenges for onset detection for them!

1. **Percussive signals:** Characterized by sharp transients, making onsets easier to detect but susceptible to false positives from noise.
2. **Harmonic signals:** Gradual onsets with stable frequencies make detection challenging, requiring more sophisticated spectral analysis.
3. **Polyphonic signals:** Multiple simultaneous onsets or overlapping events increase complexity, necessitating advanced methods to distinguish individual onsets.

2.4 Onset Detection - III. Peak Picking

2.4.1 What is the goal of peak picking in the context of onset detection?

The goal of peak picking is to identify the exact positions of local maxima in the onset detection function that correspond to actual sound onsets.

2.4.2 What are the challenges of peak picking in the context of onset detection?

Challenges include:

- Differentiating true peaks from noise-induced artifacts.
- Managing closely spaced onsets in fast sequences.
- Handling variable peak prominence due to dynamic changes in signal energy.

2.4.3 What are the processing steps for the peak picking method we discussed in the lecture?

1. **Post-processing:** Smooth the onset detection function to reduce noise and fluctuations. Apply also Normalization and DC (offset) removal.
2. Detect Local Maxima, using a **moving maximum filter**
3. Apply a threshold to filter out low-magnitude peaks.
4. **Apply adaptative thresholding** for detecting peaks to refine onset timing and eliminate spurious detections.

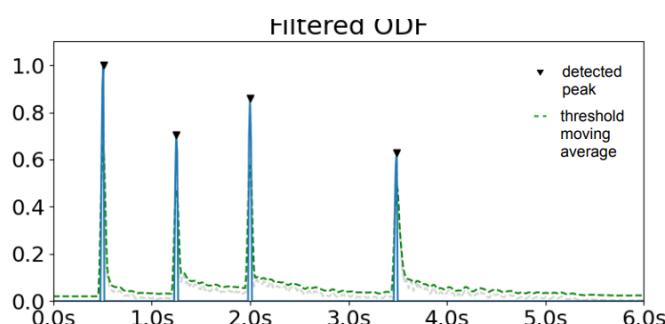


Figure 21: Post Processed ODF example

2.5 Performance Evaluation

2.5.1 What kind of annotations do we need to evaluate an onset detection method?

To evaluate an onset detection method, we need **ground-truth annotations** that specify the exact timing of each onset in the audio signal. These annotations are typically created manually by experts or derived from trusted sources.

2.5.2 What is a typical tolerance window for onset detection evaluation?

A typical tolerance window is ± 50 ms around the ground-truth onset. This accounts for slight variations in annotation precision and the temporal resolution of the detection method.

2.5.3 What evaluation metrics can be used for onset detection?

- **Precision:** The ratio of correctly detected onsets to all detected onsets.
- **Recall:** The ratio of correctly detected onsets to the total number of ground-truth onsets.
- **F-measure:** The harmonic mean of precision and recall, providing a single metric for overall performance.

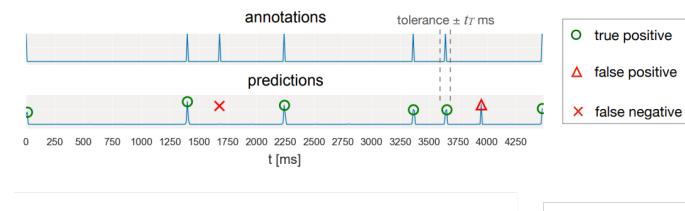


Figure 22: Example of evaluation

2.5.4 What are true positives, false positives, and false negatives in the context of onset detection evaluation?

- **True positives (TP):** Onsets correctly detected within the tolerance window of ground-truth annotations.
- **False positives (FP):** Detected onsets that do not correspond to any ground-truth onset within the tolerance window.
- **False negatives (FN):** Ground-truth onsets that are not detected by the method.

2.5.5 What are true negatives, and why are they not used for onset detection evaluation?

True negatives (TN): Instances where no onset is present, and none is detected.

They are not used in onset detection evaluation because the number of true negatives is overwhelmingly large compared to other outcomes, making it irrelevant to precision, recall, and F-measure calculations.

2.5.6 Explain precision, recall, and f-measure. How are they calculated?

- **Precision:** Measures the proportion of detected onsets that are correct.

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

- **Recall:** Measures the proportion of ground-truth onsets that are correctly detected.

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

- **F-measure:** Combines precision and recall into a single score.

$$\text{F-measure} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

2.6 State of the Art Examples

2.6.1 Give an example of a concrete implementation for an onset detector!

An example of a concrete implementation for an onset detector is **LibROSA's onset detection function**, which utilizes spectral flux as the detection function. LibROSA provides various options for preprocessing, including filtering and logarithmic scaling, as well as peak picking strategies.

A state-of-the-art approach based on spectral flux is the **LogFiltSpecFlux** method. It combines psycho-acoustically motivated preprocessing, spectral flux-based detection, and optimized peak picking to achieve high performance, particularly for certain signals.

- **Preprocessing:** Computes the STFT (window size 23ms, hop size 10ms), applies a filter bank to map the STFT to a semitone scale, and logarithmically scales the filtered magnitude spectrogram to compress the dynamic range:

$$LFSf[t] = \log(1 + \lambda \cdot FSf[t])$$

where $FSf[t]$ is the semitone-filtered spectrogram, and λ is a scaling parameter.

- **Detection Function:** Uses half-wave rectified differences to compute the spectral flux of the logarithmically scaled spectrogram:

$$dLFSF[t] = \sum_{f=0}^{F-1} \max(LFSf[t] - LFSf[t-1], 0)$$

- **Peak Picking:** A point in the detection function is considered an onset if it satisfies:

- It is a local maximum in a window (w_1, w_2) .
- It exceeds the mean in a surrounding window (w_3, w_4) .
- It maintains a minimum temporal distance (w_5) from other detected onsets.

The **LogFiltSpecFlux** method demonstrates state-of-the-art performance for onset detection in certain signal types, making it a robust and reliable implementation for audio analysis tasks.

2.6.2 Give an example for a vibrato compensation method for onset detection!

A vibrato compensation method addresses the challenge of vibrato (modulation in frequency) that can cause false positives in onset detection. One effective approach is the **SuperFlux** method, an extension of the "LogFiltSpecFlux." This method applies a **maximum filter** (± 1 quarter tones) along the frequency axis before computing the spectral difference. By compensating for frequency modulations, SuperFlux reduces false detections caused by vibrato.

Another approach is applying a **moving average filter** to smooth frequency fluctuations caused by vibrato, further minimizing the impact of these periodic variations on onset detection.

2.6.3 What is the difference between offline, online, and real-time processing?

- **Offline processing:** Analyzes the entire signal after it has been recorded, allowing access to future information for more accurate detection.
- **Online processing:** Processes the signal incrementally as it is received, without access to future data, but with some latency.
- **Real-time processing:** Processes the signal with minimal latency, suitable for live applications like live performances or interactive systems.

2.6.4 What difficulties bring online processing in the context of onset detection?

Online processing challenges include:

- Lack of future context, making it harder to confirm onsets with high confidence.
- The need for efficient algorithms to ensure low latency.
- Increased susceptibility to noise and false positives due to real-time constraints.

2.7 Onset Detection via ML

2.7.1 Why is machine learning a suitable technique for onset detection?

Machine learning is suitable for onset detection because it can learn complex patterns and features directly from data, including subtle variations in amplitude, frequency, and phase that traditional methods might overlook.

2.7.2 Describe a typical ML pipeline for onset detection!

- Preprocess audio data to create input representations (e.g., spectrograms or MFCCs).
- Annotate data with onset timings for supervised learning.
- Train a machine learning model, such as a convolutional neural network (CNN) or recurrent neural network (RNN), using the annotated data.
- Evaluate the model's performance using metrics like precision, recall, and F-measure.
- Deploy the trained model for onset detection on unseen audio signals.

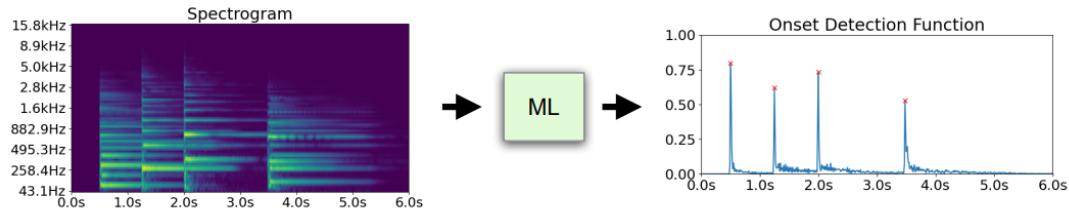


Figure 23: ML Pipeline diagram

2.7.3 How can we represent audio input for CNNs?

Audio input for CNNs can be represented as:

- **Spectrograms:** Two-dimensional time-frequency representations.
- **Mel-spectrograms:** Spectrograms transformed using the Mel scale.
- **MFCCs:** Mel-frequency cepstral coefficients for compact, perceptually relevant features.

2.7.4 What are the advantages of using ML for onset detection?

- Ability to learn complex, non-linear patterns from data.
- Adaptability to diverse audio content without manual feature engineering.
- Improved performance on challenging cases like polyphonic signals or gradual onsets.

2.7.5 What are the drawbacks of using ML for onset detection?

- Requires large, annotated datasets for effective training.
- High computational cost for training and inference.
- Potential lack of interpretability compared to traditional methods.

3 Beat Tracking

3.1 Introduction

3.1.1 What are beats, downbeats, and measures, in the context of music?

- **Beats** represent the most prominent and regular pulses in music, often aligned with foot-tapping or clapping.
- **Downbeats** are the first beats in a measure, marking the start of a bar and often emphasizing rhythmic patterns and harmonic changes.
- **Measures**, also called bars, are recurring time intervals in music that group beats according to the meter.



Figure 24: Musical Notation

3.1.2 What is tempo (in the context of music) and how can we quantify it?

Tempo is the speed at which beats occur in a musical piece. It is quantified in beats per minute (BPM), calculated as:

$$\text{BPM} = \frac{60}{\Delta t}$$

where Δt is the time interval between consecutive beats.

3.1.3 Explain the difference between local and global tempo of a music track.

Local tempo refers to the tempo of a specific section or moment in the music, which may vary significantly. Global tempo represents the average tempo of the entire track, assuming a constant or dominant tempo throughout the piece.

3.1.4 What is the goal of beat tracking?

The goal of beat tracking is to **determine the precise timing of beats in a musical signal**, aligning them with the music's temporal structure and enabling tasks such as tempo estimation, rhythm analysis, and music synchronization.

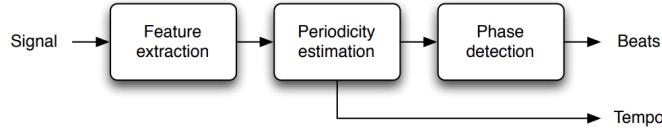


Figure 25: Pipeline

3.2 Tempo Estimation

3.2.1 What is the goal of tempo estimation?

The goal of tempo estimation is to determine the **periodicity of beats within a musical signal**, typically expressed in BPM, **to quantify the speed** of the music.

3.2.2 What are inter-onset-interval histograms, and how can they be used for tempo estimation?

Inter-onset-interval (IOI) histograms are graphical representations of the time intervals between consecutive onsets in a musical signal. Peaks in the histogram correspond to dominant periodicities, which can be mapped to tempo values in BPM to estimate the tempo.

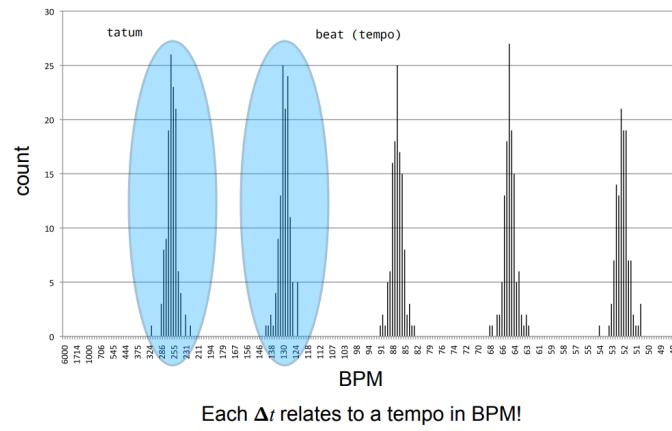


Figure 26: IOI example

3.2.3 What is autocorrelation, and how can it be used for tempo estimation?

Autocorrelation calculates the similarity of a signal with itself at varying time lags. It identifies repeating patterns in the signal, such as periodicities in an onset detection function, which can directly indicate the tempo by associating lag values with BPM.

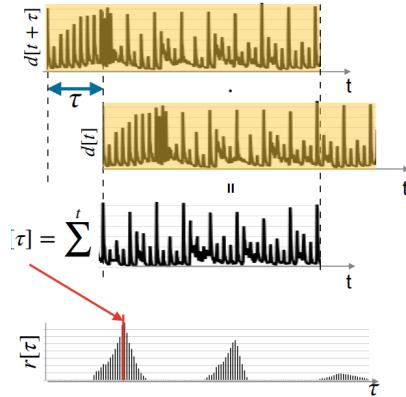


Figure 27: Autocorrelation example

3.2.4 What are octave errors in the context of tempo estimation?

Octave errors occur when the estimated tempo is an integer multiple or fraction (e.g., half or double) of the true tempo. These errors arise due to ambiguity in periodicity, where multiple tempos fit the rhythmic structure.

3.2.5 What is a tempogram?

A tempogram is a time-varying representation of tempo, calculated by analyzing periodicities in a sliding window. It shows changes in tempo over time, making it useful for tracks with significant tempo variations.

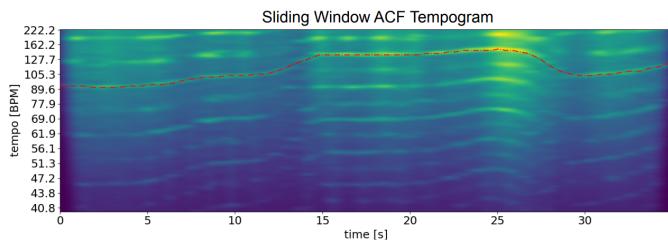


Figure 28: Tempogram example

3.3 Beat Tracking

3.3.1 What is beat phase detection, in the context of beat tracking?

Beat phase detection identifies the precise timing of individual beats within a periodic structure, aligning them with the underlying rhythm and tempo.

3.3.2 Describe a method to detect the beat phase for beat picking.

Beat phase detection can be performed using two methods:

1. **Cross-Correlation Method:** This approach uses the **onset detection function (ODF)** and a **pulse train generated from a tempo hypothesis**. The cross-correlation of the ODF with

the pulse train identifies the lag corresponding to the maximum correlation, which indicates the beat phase.

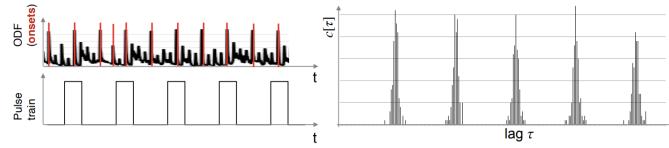


Figure 29: Pulse Train Cross-Correlation

2. **ACF Beat Tracker:** This method employs the **autocorrelation function** (ACF) of a median-filtered ODF. The filtering emphasizes salient events and reduces noise, enhancing periodicity estimation. Beats are aligned block-wise, which accommodates minor tempo variations but may struggle with large tempo changes or noisy signals. The performance of this method relies on multiple preprocessing steps, following a "weakest link" principle.

3.3.3 Explain the greedy beat tracking algorithm discussed in the lecture.

The greedy algorithm initializes the beat phase based on maximum cross-correlation. It iteratively advances by one beat period, aligning subsequent beats to the nearest onsets within a maximum allowed distance. If no onset is found, the algorithm approximates the beat position.

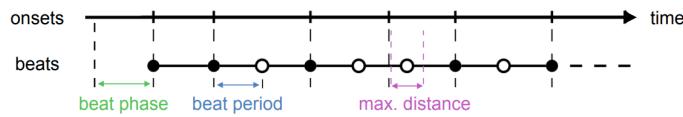


Figure 30: Greedy beat tracking

3.3.4 Explain the multi-agent tracking approach discussed in the lecture.

The multi-agent beat tracking approach generates agents corresponding to various tempo and phase hypotheses, leveraging IOI histograms to initialize tempo candidates. Each agent tracks beats independently, using a greedy algorithm to align beats with detected onsets. New agents spawn when events do not align with the current hypotheses, while similar agents can merge to consolidate tracking. A selective onset detector with high precision enhances performance. After tracking, surviving agents are scored based on beat fitness, and the beats from the highest-scoring agent are selected as the final sequence.

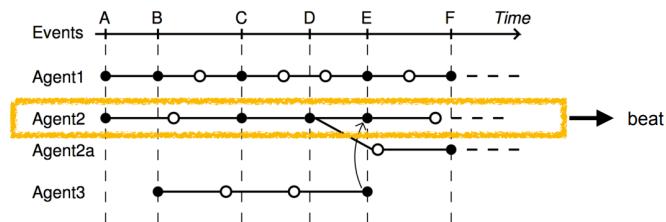


Figure 31: Multi-Agent Beat Tracker

3.3.5 Explain the machine learning setup for beat tracking.

In ML-based beat tracking, audio signals are converted into features (e.g., spectrograms) and labeled with beats and downbeats. A neural network (e.g., RNN) is trained to classify frames as beats, downbeats, or non-beats. The output probabilities are post-processed, often using a dynamic Bayesian network (DBN), to refine the predictions.

3.3.6 Why are RNNs a suitable architectural choice for beat tracking?

RNNs are well-suited for beat tracking due to their ability to model temporal dependencies in sequential data. Bidirectional RNNs, in particular, can consider both past and future context, making them effective for capturing rhythmic patterns.

3.3.7 What is the task of the RNN in the RNN beat tracker discussed in the lecture?

The RNN predicts framewise probabilities for beats, downbeats, and non-beats using temporal dependencies in input features, such as spectrograms.

3.3.8 What is the task of the DBN in the RNN beat tracker discussed in the lecture?

The DBN refines the RNN's output by jointly inferring tempo, meter, and phase, ensuring consistency across beats and downbeats while allowing tempo changes.

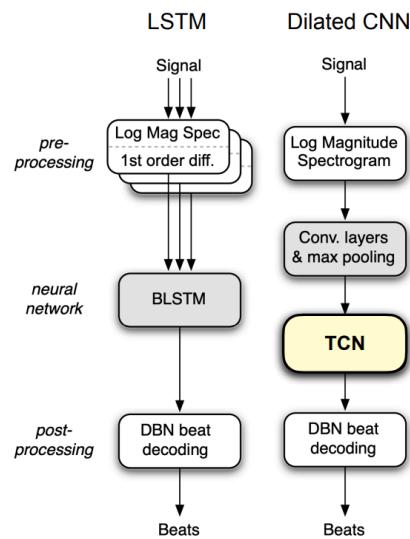


Figure 32: LSTM & Dilated CNN architectures

3.4 Evaluation

3.4.1 Which metrics can we use for beat tracking evaluation?

Metrics for evaluation include:

- **F-measure:** Measures precision and recall within a ± 70 ms tolerance.

- **P-score:** Evaluates tracking accuracy based on deviations within 20% of the annotated beat interval.
- **CMLc and CMLt:** measure the longest continuously segment (CMLc) or all correctly tracked beats (CMLt) at the correct metrical level. A beat is considered correct if it is reported within a 17.5% tempo and phase tolerance, and the same applies for the previously detected beat.
- **AMLc and AMLt:** like CMLc & CMLt, but additionally allow offbeat and double/half as well as triple/third tempo variations of the annotated beats.

3.4.2 In what range are tolerances for beat positions typically?

Tolerances for beat tracking are typically in the range of **±70ms**, depending on the application.

3.4.3 Why are tolerances for beat tracking higher than for onset detection?

Beat tracking tolerances are higher because beats may not align perfectly with onsets. Rhythmic structure and musical context allow for flexibility, making exact timing less critical than for onset detection.

4 F0 and Notes

4.1 Pitch and F0

4.1.1 Properties of Musical Notes Mentioned in the Lecture

- **Onset Time:** The time when a note begins.
- **Offset Time:** The time when a note ends.
- **Fundamental Frequency (f_0):** The pitch or base frequency of a note.
- **Relative Loudness (Velocity):** The volume or intensity of the note.
- **Instrument (Timbre):** The unique quality or color of the sound.
- **Playing Techniques:** Variations such as glissando, vibrato, tremolo, harmonic/percussive components, etc.

4.1.2 Pitch and Its Relationship to the Fundamental Frequency of a Note

- **Pitch:** A perceptual quantity that allows humans to order sounds on a frequency-related scale from low to high.
- **Relationship:** Pitch perception is closely tied to the fundamental frequency (f_0) of a sound, which is often derived logarithmically.

4.1.3 Harmonics or Harmonic Partials of a Note

- **Harmonics:** Integer multiples of the fundamental frequency (f_0).
- **Harmonic Partials:** Components of a sound that resonate at frequencies higher than f_0 , contributing to the timbre of the note.

4.1.4 Fundamental Frequency of a Note

- The **fundamental frequency (f_0)** is the lowest frequency of a sound and is responsible for the perceived pitch of the note.

4.1.5 Difference Between Percussive, Harmonic, and Harmonic-Percussive Notes

- **Percussive Notes:** Do not have a distinct fundamental frequency; their sound arises from transient noise (e.g., drums).
- **Harmonic Notes:** Have a clear fundamental frequency and harmonics (e.g., violin or flute).
- **Harmonic-Percussive Notes:** Combine properties of both percussive and harmonic sounds (e.g., piano).

4.1.6 Visualization Method for Monophonic F0 Estimation/Annotations

The f_0 **Contour:** A common visualization method for tracking f_0 over time for monophonic melodies, often calculated from spectrograms.

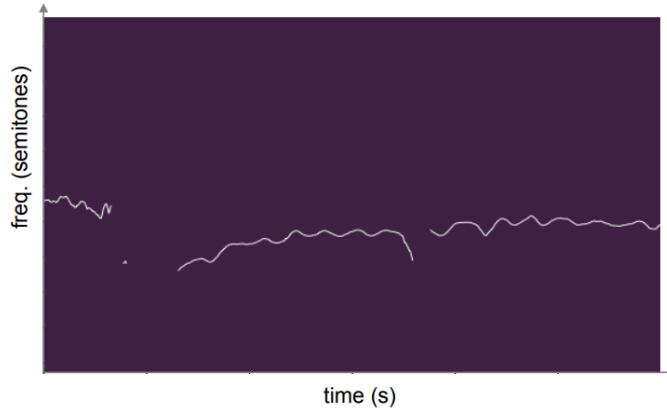


Figure 33: F0 contour

4.1.7 Visualization Methods for Polyphonic F0 Estimation/Annotations

The **Piano Roll Representation**: Suitable for polyphonic music, often used in digital audio workstations. It is akin to a spectrogram but without harmonics and transients.

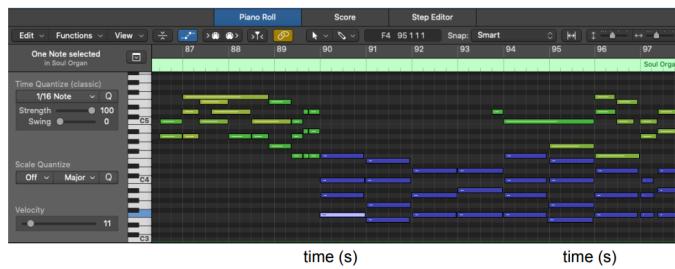


Figure 34: Piano Roll representation

4.1.8 Difference Between Single- and Multi-F0 Estimation

- **Single-F0 Estimation:**
 - Detects the f_0 of monophonic signals (e.g., singing voice).
 - Accurate and robust algorithms exist; considered a solved problem.
- **Multi-F0 Estimation:**
 - Detects the f_0 of each tone in polyphonic signals (e.g., piano or mixed instruments).
 - Still an open research problem.

4.1.9 Three Levels for F0 Estimation Discussed in the Lecture

1. **Frame-Level Detection:** Focused on detecting pitch contours. We estimate f_0 for each frame
2. **Note-Level Detection:** Used for monophonic transcription. We need to do a post-processing for framewise detection. There is a simultaneous onset/pitch/offset detection.

3. **Stream-Level Detection:** Applied in automatic music transcription. There is a post-processing for note detection and instrument note detection

4.2 Evaluation

4.2.1 Evaluation for F0 Estimation on Frame-Level

For F_0 estimation on the frame level, annotations are discretized to a frame rate. This results in a continuous F_0 contour that is used for evaluation.

- Each frame is evaluated individually.
- Each sounding pitch is evaluated separately.
- An absolute pitch difference of less than 50 cents (i.e., a quarter tone) is typically considered correct.
- Common metrics include Precision, Recall, F-measure, or Accuracy.

4.2.2 Evaluation for F0 Estimation on Note-Level

Note-level evaluation involves a piano-roll representation, where annotations are represented as **discrete quantized** notes. This method emphasizes musical properties such as the onset and offset times of the notes.

- Solves some of the problems associated with frame-wise evaluation.
- Suitable only for discrete, quantized notes (not continuous pitch contours).
- Various aspects of notes can be evaluated, such as:
 - Onset (e.g., within $\pm 50\text{ms}$) and pitch (e.g., within ± 50 cents).
 - Onset, pitch, and offset (e.g., max $\pm 50\text{ms}$, 20% of note length).
- Metrics are usually measured independently, without an overall combined metric.

4.2.3 Evaluation for F0 Estimation on Stream-Level

Stream-level evaluation treats each instrument separately at the note level. For multi-instrument setups, each instrument is evaluated independently on its note-level representation.

4.2.4 Issues with Frame-Level F0 Estimation

Frame-level evaluation is convenient but problematic:

- Short missing sounding frames have marginal impact on overall performance but are very irritating to humans.
- Similarly, falsely detected short spurious notes are problematic in frame-level evaluation.

4.2.5 Metrics for Frame-Level F0 Estimation

The metrics commonly used for frame-level F_0 estimation include:

- Precision
- Recall
- F-measure
- Accuracy: $\text{acc} = \frac{TP+TN}{TP+FP+TN+FN}$, often assuming $TN = 0$.

4.2.6 Issues with Note-Level F0 Estimation

Note-level F_0 estimation has its own challenges:

- Determining accurate onset and offset times can be subjective and unclear from signal information alone.
- Errors in note segmentation or classification significantly affect the results.

4.2.7 Properties Evaluated for Note-Level F0 Estimation

The properties typically evaluated for note-level F_0 estimation include:

- Note onset time
- Note offset time
- Pitch accuracy
- Note duration

Various aspects of notes can be evaluated

- Onset (e.g. $\pm 50\text{ms}$) + pitch (e.g. $\pm 50\text{cents}$)
- Onset + pitch + offset (e.g. max $\pm 50\text{ms}$, 20% of note length)

Usually measured independently (no “overall” metrics)

4.3 Fundamental Frequency Estimation

4.3.1 Which simple/naive methods could be used for F0 estimation?

- **Zero-Crossing Rate (ZCR):** Measures how often the waveform crosses zero per unit of time. This was one of the first approaches to F_0 -estimation. ZCR is directly related to the number of times the waveform repeats per unit time.
- **Autocorrelation:** Analyzes periodicity by finding the lag at which the signal best correlates with itself. Often used for pitch detection in periodic signals.

4.3.2 Explain the typical problems of simple F0 estimation methods!

- **Problems with ZCR:**
 - ZCR works only if the signal does not contain extra zero-crossings.

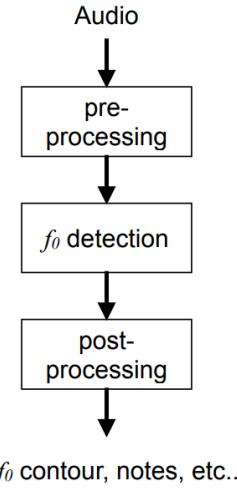


Figure 35: F0 estimation pipeline

- Partial-rich tones produce too many zero-crossings, leading to incorrect frequency estimation.
- Estimated frequency can be off by a factor of 2 if there are 4 zero-crossings instead of 2 per cycle.
- Requires preprocessing, such as low-pass filtering, to counteract the effects of additional crossings.

- **Problems with Autocorrelation:**

- Partial-rich tones may produce a dominant peak at one of the harmonic partials instead of the fundamental frequency.
- When the signal is pseudo-periodic with a low-power fundamental frequency, even humans may mistake a harmonic partial for the fundamental.

- **General Issues with Simple Methods:**

- Most common errors are **octave errors**, where the first partial is detected instead of the fundamental ($f = 2 \cdot f_0$).
- Naive methods struggle to cope with harmonic-rich signals or noise.

4.3.3 Explain the processing steps for the YIN F_0 estimation method!

1. **Difference Function Calculation:**

- YIN operates in the time-domain, similar to the Autocorrelation Function (ACF) method.
- Replaces ACF with the Average Magnitude Difference Function (AMDF):

$$d_t[\tau] = \sum_{t=1}^M (x[t] - x[t + \tau])^2$$

- Utilizes the principle of cancellation: $d_t[T] \approx 0$ for periodic signals with period T .

2. Normalization:

- Values of $d_t[\tau]$ depend on the range of $x[t]$, so a simple minimum search is insufficient.
- Normalization produces the Cumulative Mean Normalized Difference Function $nd_t[\tau]$:

$$nd_t[\tau] = \begin{cases} 1 & \text{if } \tau = 0, \\ \frac{d_t[\tau]}{\frac{1}{\tau} \sum_{i=1}^{\tau} d_t[i]} & \text{otherwise.} \end{cases}$$

- This reduces valleys caused by strong first partials and ensures the range of $nd_t[\tau]$ lies within $[0, 2]$, where 0 indicates perfect periodicity.

3. Thresholding:

- After normalization, an absolute threshold is applied to identify candidates.
- The smallest lag τ below the threshold is chosen as the fundamental frequency.
- If no values fall below the threshold, the global minimum is used as a fallback.

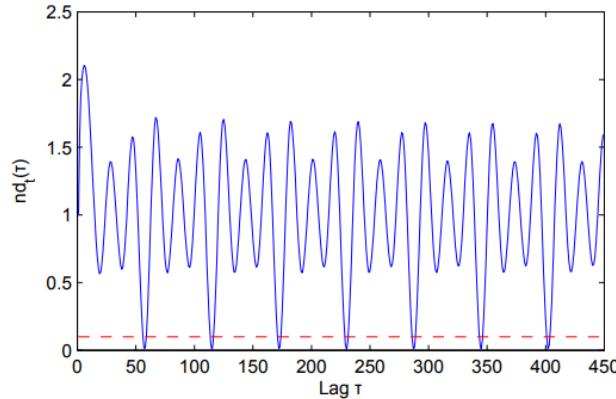


Figure 36: Thresholding in YIN

4. Interpolation:

- To address quantization errors, YIN employs parabolic interpolation.
- Uses $nd_t[\tau - 1]$, $nd_t[\tau]$, and $nd_t[\tau + 1]$ to fit a parabola.
- The minimum of the fitted parabola is selected as the refined lag τ .

4.3.4 What is the difference between YIN and pYIN?

- **YIN:**
 - A deterministic F_0 estimation algorithm based on the time-domain.
 - Outputs a single F_0 value or absence for each frame.
- **pYIN (Probabilistic YIN):**
 - Extends YIN by incorporating a probabilistic framework.

- Produces a probability distribution over F_0 values for each frame instead of a single output.
- Improves robustness against noisy signals and enhances pitch tracking by considering multiple candidates.

Feature	YIN	pYIN
Core Methodology	Deterministic, time-domain analysis	Probabilistic extension of YIN
Output	Single F_0 value or absence per frame	Probability distribution over F_0 values
Noise Robustness	Moderate	High
Voicing Detection	Implicit	Explicit
Continuity in Output	May have jumps or discontinuities	Smoother pitch contours
Computational Cost	Lower	Higher
Use Case	Clean, monophonic signals	Noisy, real-world, or polyphonic signals

Table 1: Comparison of YIN and pYIN

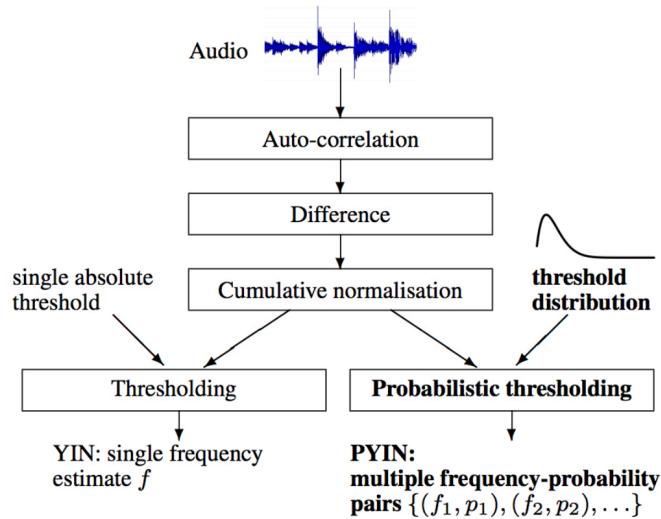


Figure 37: YIN and PYIN pipeline

4.4 F0 Estimation in the Frequency Domain

4.4.1 What is the fundamental idea behind F_0 estimation in the frequency domain?

The fundamental idea is to analyze the spectrogram and locate the peaks corresponding to harmonic frequencies. The F_0 is estimated by identifying the fundamental frequency or by analyzing the spacing between harmonics. This method historically represents earlier approaches to pitch estimation and has the appeal of being generalizable to multiple F_0 estimation.

4.4.2 List some of the problems of F_0 estimation in the frequency domain!

- Simply picking the bin with the most energy is insufficient:
 - Harmonics may have a larger magnitude than the fundamental F_0 .
 - In some cases, the fundamental F_0 may have no energy (missing fundamental).

- Common challenges include:
 - **Harmonic errors:** $f_n \neq F_0 \cdot k$, where $k \in \mathbb{N}$.
 - Missing harmonics that disrupt the detection of F_0 .
 - Octave errors, where higher harmonics are mistaken for the fundamental.
- For multiple F_0 :
 - Overlapping harmonics and transients make detection challenging.
 - Determining how many peaks (or notes) to select can be ambiguous.

4.4.3 How does the SHAPE F_0 estimation method work, in principle?

The SHAPE (Smooth Harmonic Average Peak Envelope) method works by using spectral templates and a smooth harmonic averaging approach:

- Define a **pitch template function** representing the expected harmonic structure of a given pitch.
- Compute the correlation of this template with the actual spectrum ($|X[t]|$) over a set of candidate pitches.
- Select the pitch with the **highest correlation** as the estimated F_0 .

Additionally, SHAPE employs the **Smooth Harmonic Average Peak-to-Valley Envelope**, which:

- Calculates the average peak-to-valley distance at harmonic locations.
- Combines techniques to address common challenges:
 - Penalizes non-harmonic frequency content.
 - Applies decaying weights for higher harmonics.
 - Ensures smooth weight transitions.

This approach ensures more robust F_0 estimation by reducing the influence of noise and non-harmonic components. See Figure 38 for an example.

To address specific problems:

- Frequencies closer to F_0 are weighted more heavily.
- Penalize incorrect harmonics by considering the entire spectrum, not just harmonic peaks.
- Handle imperfect harmonics (non-integer multiples) by adapting the template to the specific signal.

This approach provides a more robust solution to issues like missing fundamentals and harmonic interference.

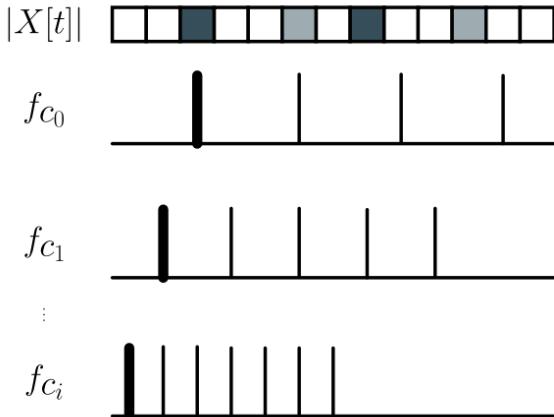


Figure 38: Example of spectral templates used in pitch estimation.

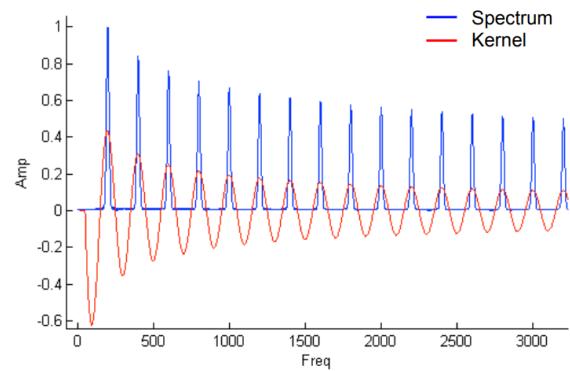


Figure 39: SHAPE example.

4.5 Multiple F0 Estimation

4.5.1 Explain the difference between single and multiple- F_0 estimation!

- **Single- F_0 Estimation:**

- The task is to estimate a single fundamental frequency (F_0) present in a monophonic audio signal.
- Assumes that there is only one source active at any given time.
- Common in applications like voice pitch detection or simple melody tracking.

- **Multiple- F_0 Estimation:**

- The task is to estimate multiple F_0 values simultaneously from a polyphonic audio signal.
- Accounts for overlapping harmonics, transients, and the presence of multiple instruments or voices.
- Common in applications like polyphonic music transcription or source separation.

4.5.2 What are the challenges of multiple- F_0 estimation?

- **Overlapping Harmonics:**

- Harmonics from different sources can overlap, making it hard to distinguish individual F_0 values.

- **Diverse Instrument Sounds:**

- Different instruments produce distinct harmonic structures and timbres, complicating harmonic identification.

- **Unknown Number of Notes:**

- It is often unknown how many simultaneous F_0 values exist in the signal.

- **Noisy Spectrogram:**

- Overlapping harmonics and fundamentals create a noisy spectrogram that is difficult to interpret.

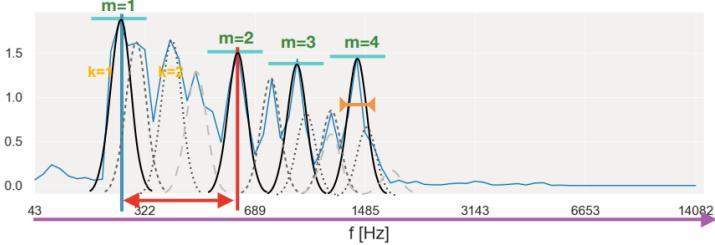
4.5.3 Explain the three categories for multiple- F_0 estimation we discussed in the lecture!

1. Spectrogram Analysis:

- Detect peaks corresponding to fundamental frequencies and harmonics.
- Challenges include overlapping harmonics and missing fundamentals.

Probabilistic Gaussian mixture model

[Goto, 2004],
 [Kameoka et al., 2007],
 [Yoshii & Goto, 2012]



Notes:

$$\mathcal{M}_k(f) = \sum_{m=1}^M \tau_{k,m} \mathcal{N}(f | \mu_k + o_m, \sigma_k)$$

frequency harmonics amplitude f_0 harmonics offset width

Spectrum:

$$\mathcal{M}_t(f) = \sum_{k=1}^K \pi_{t,k} \mathcal{M}_k(f)$$

notes

Figure 40: Spectrogram Analysis F0

2. Spectral Decomposition:

3. Decompose the spectrogram into components using techniques like:

- Non-Negative Matrix Factorization (NMF).
- Probabilistic Latent Component Analysis (PLCA).
- Sparse Coding.

These methods model the spectrogram as a combination of templates and activations.

4. Machine Learning Approaches:

- Treat F_0 estimation as a multi-label classification problem.
- Use algorithms like:
 - Support Vector Machines (SVMs).

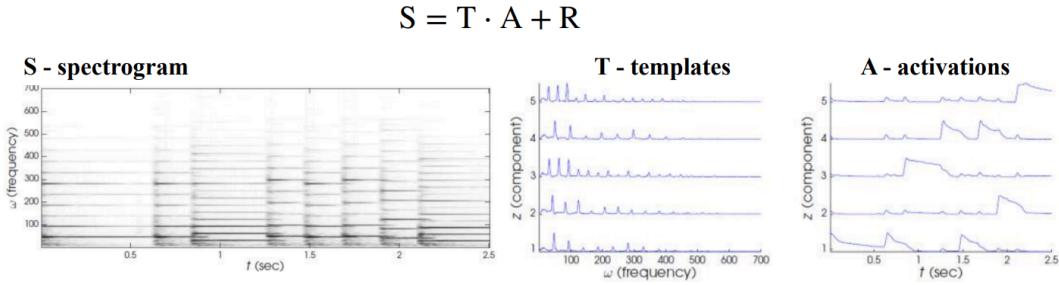


Figure 41: Spectral Decomposition in F0

- Hidden Markov Models (HMMs).
- Neural Networks (NNs).
- Requires a large amount of training data to learn from examples.

4.6 Neural Network Based Methods

4.6.1 Explain how a CNN can be used to perform multi-f0 estimation!

A Convolutional Neural Network (CNN) can be used for multi-f0 estimation by operating on spectrogram representations of polyphonic audio. The spectrogram provides a time-frequency representation of the audio, capturing harmonic structures that are crucial for pitch detection. The CNN extracts patterns corresponding to harmonic frequencies using convolutional layers. Each frame of the spectrogram is processed independently, and the network predicts the active pitches for each frame. The output consists of a binary or continuous-valued vector representing the presence of each pitch (e.g., 88 piano keys) at a given time frame. This frame-wise transcription acts as an "acoustic model," focusing on the spectral features associated with pitch activation.

4.6.2 What issues can arise when doing frame-wise multiple-f0 estimation?

- **Temporal Continuity:** Frame-wise processing lacks temporal context, leading to inconsistencies in pitch activation across frames.
- **Overlapping Harmonics:** Harmonics from different pitches may overlap, making it challenging to distinguish individual pitches.
- **Transients:** The attack and decay phases of notes can be misinterpreted as new notes or missed entirely.
- **Note Sustain:** Sustained notes may be inconsistently detected, especially if harmonics decay unevenly.

4.6.3 What is the task of the HMM post-processing for the CNN multiple-f0 estimation approach, we discussed in the lecture?

The Hidden Markov Model (HMM) post-processing aims to enforce temporal consistency in the CNN's frame-wise predictions. Its tasks include:

- **Temporal Smoothing:** Reducing abrupt transitions in pitch activations to provide a smoother representation.
- **State Constraints:** Guiding transitions between pitch states (e.g., onset, sustain, release) based on predefined rules.
- **Noise Reduction:** Suppressing spurious activations caused by noise or overlapping harmonics by modeling the note activation process.

4.6.4 Why does explicit modeling of note phases (ADSR) improve performance for multiple-f0 estimation?

Explicit modeling of note phases, such as Attack, Decay, Sustain, and Release (ADSR), improves performance by:

- **Temporal Representation:** Capturing the temporal evolution of notes ensures accurate identification of note onsets, sustains, and releases.
- **Context Awareness:** Phase-specific outputs help the network distinguish between the start, middle, and end of a note, reducing false positives and negatives.
- **Granular Training Labels:** Providing phase-specific training labels allows the network to learn the dynamics of note transitions, improving transcription accuracy.

5 Music Classification and Semantic Annotation

Recap: Audio Analysis So Far

- Extraction of features (low-level).
- Detection of events:
 - Onsets.
 - Beats.
- Description of tonal properties:
 - Pitch estimation.

5.1 Semantic Description

Semantic description refers to the use of words that make sense to humans when describing music. These descriptions focus on musical properties that hold semantic meaning for the average listener. For example, terms like "genre," "mood," and "instruments" provide intuitive ways to characterize music. Semantic representation can be categorized into three levels:

- **High-level representations:** Musical concepts as humans perceive them, such as emotions or style.
- **Mid-level representations:** Serve as a bridge, combining low-level features into meaningful patterns.
- **Low-level representations:** Consist of statistical descriptions of audio signals that are machine-readable.

5.2 Musical Genre

A musical genre is defined as a set of musical events, real or possible, governed by a set of socially accepted rules. These rules can be categorized into five types, as proposed by Fabbri in 1981:

1. **Formal and technical rules:** Content-based practices.
2. **Semiotic rules:** Abstract concepts, such as emotions or political messages conveyed through music.
3. **Behavioral rules:** Define how composers, performers, and audiences interact with the genre.
4. **Social and ideological rules:** Describe the connection between genres and demographic factors, like age or culture.
5. **Economic and juridical rules:** Focus on systems that support genres, such as contracts and performance venues.

5.3 Why Genre Classification?

Genre classification is essential for:

- Grouping songs based on common criteria and labeling musical trends.
- Organizing and facilitating the production, circulation, and reception of music.
- Supporting Music Information Retrieval (MIR) as a common evaluation method and a proxy for music similarity tasks.

Archetypal genres, such as rock or jazz, are often immediately recognizable, sometimes within just 300–400 milliseconds. MIR commonly uses genre classification for its simplicity compared to pairwise similarity testing, relying on fixed ground truths like the "Cranfield paradigm."

5.4 Genre Classification

Genre classification is modeled as a machine learning task. It involves:

- Extracting features from musical data, ranging from low-level signal attributes (e.g., spectral flatness) to more abstract descriptors (e.g., rhythm patterns).
- Training a classifier on labeled examples.
- Using the trained classifier to categorize unseen data.

5.5 Case Study: k-NN for Genre Classification

A notable case study explored the use of a k-Nearest Neighbors (k-NN) classifier for genre classification:

- Parameters:
 - $k = 3$.
 - Features used: MFCC and spectral flatness.
- Dataset: GTZAN with a 10-fold cross-validation setup.
- Results:
 - Accuracy: 80.8%.
 - Misclassified examples: 192.

5.6 Why NOT Genre Classification?

Genre classification faces several challenges:

- **Cultural differences:** Lead to varying interpretations of the same music, making universal agreement difficult.
- **Expert disagreements:** Inter-listener agreement is reported at only 80%.
- **Genre fluidity:** Intensional definitions of genres may not exist, as genres are fluid and dynamic.

- **Complexity of music:** For example, the "Bohemian Rhapsody effect," where a single label is insufficient to capture the complexity of a track, which may simultaneously fall under multiple categories (e.g., ballad, pop, rock, musical, and metal).

5.7 Tags

Tags are short descriptions of specific aspects of music, such as:

- Genre, style, instrumentation, or mood.
- Performer characteristics or personal categories (e.g., "favorite" or "seen live").

Tags can be:

- **Explicit:** Assigned by experts.
- **Implicit:** Derived from user interactions.

Weighted tags reflect the number of users who agree on a specific tag. Common methods for obtaining tags include:

- Expert annotations.
- Collaborative tagging.
- Games designed for tagging purposes.

5.8 Games With A Purpose (GWAP)

An innovative approach to manual tagging involves using games to engage users while collecting meaningful data. Known as Games With A Purpose (GWAP), these tasks integrate tagging into gameplay, encouraging users to put effort into completing the task.

5.8.1 ESP Game

The ESP Game employs a collaborative approach:

- Simulates players by replaying sessions.
- Includes features like "taboo words," restricting commonly used terms to encourage diverse input.
- Removes overused or skipped images to ensure a wide variety of content is tagged.

This method enhances the quality and breadth of data collected.

5.8.2 TagATune

TagATune focuses on "input agreement" rather than "output agreement." Participants describe a tune to determine whether they are listening to the same or a different one as their partner. Advantages include:

- Avoiding semantic gaps commonly found in output-based systems.
- Collecting high-quality tag data.

However, this game revealed that output agreement is less effective for music, highlighting the challenges in bridging the semantic gap.

5.8.3 MagnaTagATune Dataset

The MagnaTagATune dataset is a product of the TagATune game. Key features include:

- Over 25,000 snippets from more than 5,000 songs.
- Detailed annotations, including tag-based descriptions, similarity matrices, and expert-provided genre labels.

This dataset has been widely used for tasks like music recommendation and classification.

5.9 Audio Feature Extraction

5.9.1 Mel Frequency Cepstral Coefficients (MFCCs)

MFCCs represent the spectral envelope of audio, capturing timbral aspects critical for sound similarity. Steps for computation:

1. Framing the signal.
2. Applying the Discrete Fourier Transform (DFT).
3. Mapping the spectrum to the Mel scale for perceptual relevance.
4. Taking the logarithm of Mel-scaled amplitudes to reflect perceived loudness.
5. Performing the Discrete Cosine Transform (DCT) to decorrelate features.

5.9.2 The Mel Scale

The Mel scale maps frequencies in Hertz to their corresponding Mel values using the formula:

$$m = 2595 \log_{10}\left(1 + \frac{f}{700}\right)$$

Typically, 40 Mel filter bins are used in MFCC computation.

5.9.3 Bag-of-Frames Modeling

In Bag-of-Frames modeling:

- An audio track is represented as a set of MFCC vectors.
- These vectors are aggregated using statistical measures like mean, variance, or covariance.
- Comparisons are made using distances between feature distributions, e.g., Earth Mover's Distance or KL Divergence.

While effective, this approach loses temporal information.

KL Divergence for MFCCs

KL Divergence measures the difference between probability distributions. For Gaussian models of MFCCs, a closed-form solution exists, enabling efficient comparison. Applications include using MFCC "Bag-of-Frames" with Support Vector Machines (SVMs) for genre classification.

5.10 Deep Learning Approaches

Deep learning enables end-to-end learning by directly extracting features from raw signals. Key features include:

- Spectrograms or mel-spectrograms as input.
- Architectures like Convolutional Neural Networks (CNNs), Deep CNNs, and CRNNs.

These models predict single or multi-class labels (e.g., genres or tags). Careful design is necessary to:

- Handle percussive versus harmonic properties.
- Ensure pitch invariances.

5.11 Deep Learning for Music Tagging and Classification

5.11.1 End-to-End Learning for Tags

End-to-end learning utilizes CNNs to:

- Automatically extract audio features.
- Ensure translation, distortion, and locality invariance.

The architecture processes log-amplitude mel-spectrograms, using:

- Multiple convolutional layers with activation functions like ReLU.
- Batch normalization and dropout for regularization.
- A sigmoid activation function for predicting multiple tags.

5.11.2 CNN Architectures and Design

The architecture for tagging involves:

- Input: 29.1-second audio clips downsampled to 12 kHz, resulting in 1,366 frames per clip.
- Features: Log-amplitude mel-spectrograms with 96 mel bands.
- Processing: Multiple convolutional layers with varying filter sizes (e.g., 3x3) and max pooling layers for downsampling.
- Optimization: Batch normalization, dropout, and ADAM optimizer.
- Output: 50 tags per clip using a sigmoid activation function.

5.11.3 Comparison of Input Representations and CNN Models

CNNs provide competitive results for music tagging:

- Lightweight and easy to train.
- Often involve many layers or large filters to achieve a wide receptive field.
- Dilated convolutions enhance the receptive field by skipping pixels, making tagging reliant on global properties.

Experiments reveal:

- Preprocessing methods include raw waveform (PCM), STFT, mel-spectrogram, CQT, and MFCCs.
- Evaluations performed on datasets like MagnaTagATune and MTG-Jamendo.
- Dilated CNNs perform better on two-dimensional input representations like mel-spectrograms but struggle with raw waveforms.
- Statistical tests, such as two-way ANOVA, confirm significant influences of model architecture and input representation.

5.11.4 Results from Choi et al. (2016)

- The proposed system (FCN-4) achieved an AUC of 0.894 on the MagnaTagATune dataset.
- End-to-end CNN models outperform handcrafted feature-based systems.
- Dilated CNNs excel in training efficiency and produce robust results on diverse input formats.

5.11.5 Experimental Setup and Analysis

Five preprocessing methods were evaluated using corresponding CNN models, including VGG-16, ResNet, SENet, MusiCNN, and Dilated CNN:

- **Datasets:** MagnaTagATune (25,000 samples) and MTG-Jamendo (55,000 samples).
- **Metrics:** ROC-AUC and PR-AUC.
- **Statistical Analysis:** Multiple runs per configuration with two-way ANOVA and Tukey-HSD tests.

Results showed:

- STFT preprocessing achieved the best performance, while raw waveforms performed the worst.
- Lightweight MFCC-based approaches yielded consistent results.

6 Multimodality

Music is represented in various formats across multiple domains, impacting how it is perceived. It is not merely treated as a signal but as a cultural phenomenon. This perspective allows for deeper analytical and cultural explorations. Semantic tagging and description approaches are heavily influenced by the availability and quality of data, enabling richer interpretations and applications.

6.1 Diverse Music Data Sources

Music data comes from a variety of sources:

- **Content-based:** Includes audio, symbolic data, and lyrics. Features can be extracted directly from audio files, enabling recommendation systems without reliance on community data or cultural biases.
- **Meta-data:** Editorial and curatorial data, as well as multi-modal sources like album covers, enrich metadata-driven recommendations.
- **User-generated data:** Tags, reviews, blogs, biographies, and personal stories allow for personalized insights.
- **Interaction data:** Logs, feedback mechanisms (e.g., thumbs), and purchase history help model user preferences.
- **Curated collections:** Playlists, radio channels, and album compilations provide structured user preferences.

6.2 Collaborative Filtering (CF)

Collaborative Filtering leverages user interaction data to recommend music based on preferences:

- CF relies on implicit data (e.g., plays) and explicit data (e.g., likes/dislikes).
- The task involves completing a sparse user-item matrix by predicting unknown interactions.
- It assumes users with similar tastes in the past will exhibit similar preferences in the future.

The two main methods are:

1. *User-based CF:* Relies on the similarity of users.
2. *Item-based CF:* Relies on the similarity of items.

6.3 Factors Hidden in the Data

Matrix factorization-based recommender systems assume that observed data (e.g., ratings) represent interactions between users and items. Latent factors extracted through matrix decomposition represent underlying patterns, aiding in user-item interaction predictions.

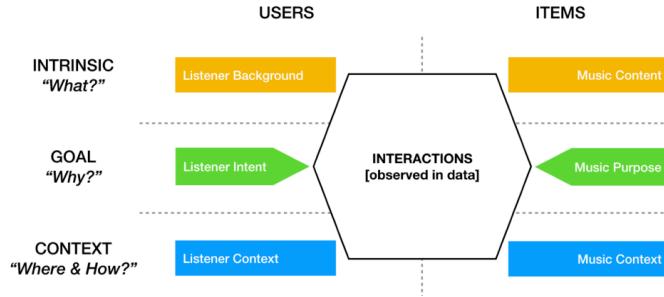


Figure 42: Factors Hidden in Data

6.4 Audio Content Analysis

Audio-based recommendation systems provide true content-based recommendations by extracting features directly from audio files, removing the need for community or cultural data:

- High-level semantic descriptors are learned from low-level features using machine learning.
- Representation learning and temporal modeling are state-of-the-art approaches.

Advantages of this approach include:

- No cultural biases.
- No popularity biases.
- Independence from subjective ratings.

6.5 Recap: Audio Content Analysis (2017 State-of-the-Art)

Audio content analysis involves:

- **Timbre:** Features like MFCCs are used for genre classification and “more-of-this” recommendations.
- **Beat/Downbeat:** Tempo extraction (e.g., using madmom for 85 bpm).
- **Tonal Features:** Pitch-class profiles for melody extraction or cover version identification (e.g., Essentia).
- **Semantic Categories:** Machine learning enables predictions like *not_danceable*, *mood_not_happy*, etc.



Figure 43: Target in Audio Content Analysis

6.6 Text Analysis Methods (Basic IR)

Text analysis leverages user-generated content and lyrics to capture aspects beyond the audio signal. Methods include:

- Bag-of-Words, Vector Space Models, TFIDF.
- Advanced methods like topic modeling, word2vec, BERT, and representation learning.

Applications include:

- Tag-based similarity.
- Sentiment analysis.
- Mood detection in lyrics.



Figure 44: Target in text analysis

6.7 Textual Sources of Music Similarity

Text processing of user-generated data (e.g., reviews, blogs, tags) is critical for music similarity. However:

- The process involves noisy and potentially erroneous data.
- Similarity measures (e.g., cosine distance) are applied.
- Personalization remains limited as it depends on user-generated content.

6.8 Multimodal Approaches

Multimodal approaches integrate multiple sources of information to enhance recommendation systems, particularly to address cold-start problems. Examples include:

- Combining text embeddings (e.g., artist biographies) with CNN-trained audio embeddings.
- Fusing deep features from audio, image (e.g., album covers), and text.



Figure 45: Target in multimodal approaches

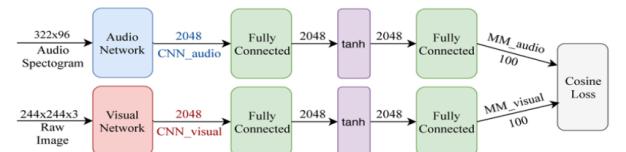


Figure 46: Example of multimodal analysis

6.9 Feedback-Transformed Content

Content features are aligned with collaborative filtering (CF) models to address cold-start problems. Approaches include:

- Predicting CF data from audio features.
- Personalizing mixtures of content features.
- Learning item-based CF similarity functions using metric learning.
- Weighted matrix factorization with CNN inputs (e.g., mel-spectrograms).
- Tag-trained neural networks integrated with matrix factorization to emphasize content.

6.10 Beyond Items

Describing items through non-personalized models helps create generic recommendations. These models use aggregated user data to focus on:

- Cultural or contextual metadata.
- Aspects independent of individual users.

6.11 Listener Background

Psychological and sociological factors are vital for building predictive user models and addressing cold-start issues. Factors include:

- User personality, music preferences, demographics, and cultural context (e.g., age, mood, individualist vs. collectivist cultures).
- Findings:
 - Preferences vary by age, mood, and cultural background.
 - Introverts may prefer sad music, while open individuals may seek uplifting music when sad.

6.12 Listener Context

The user's context significantly affects music recommendations. Types of context include:

- **Environment-related:** Factors like time, location, weather, and traffic conditions.
- **User-related:** Includes activity, emotion, and cultural context.

Context data is obtained through:

1. **Explicit methods:** User input via questionnaires.
2. **Implicit methods:** Sensors in smart devices like accelerometers or noise detectors.
3. **Inference:** Using rules or machine learning techniques (e.g., inferring mood from skipping behavior or activity from movement data).

6.13 Services, Data Bias and Limitations

Music recommendation systems and models rely heavily on the data used for training and evaluation. However, significant limitations and biases exist in this data:

- **Source of Data:** The data used in training listening and preference models often originates from specific platforms or services, meaning there is no true "listening in the wild."
- **Platform/Service Bias:** The platform or service acting as a mediator impacts the type of data generated. This can influence listening behaviors and preferences recorded in the dataset.
- **Lack of Generalizability:** Data collected from one service is not necessarily generalizable to other platforms or contexts.
- **Data Variability:** Substantial differences exist in data quality and type across services, further complicating the creation of unbiased models.

These factors highlight the challenges of developing universal and unbiased recommendation systems, as the quality and diversity of data fundamentally shape the models' effectiveness and applicability.

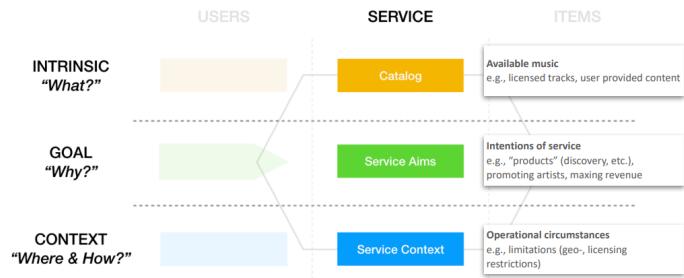


Figure 47: Services in multimodality

7 Drum Transcription

7.1 Introduction

7.1.1 Explain the automatic drum transcription (ADT) task!

A **drum** is a percussion instrument that produces sound by striking a stretched membrane or surface, typically with sticks, hands, or mallets. The sound is generated through the vibration of the drumhead, which is stretched over a hollow body called the shell. Drums come in various types, such as the bass drum (low, deep tones), snare drum (sharp, staccato sounds), and hi-hat (used for rhythmic patterns). These instruments form the foundation of rhythm and timing in music, making their transcription essential for understanding and recreating drum patterns in audio.

The automatic drum transcription (ADT) task aims to analyze audio signals to identify and transcribe **drum events**, including their timing and instrument classification (e.g., bass drum, snare drum, hi-hat). It enables applications such as **generating sheet music, remixing, and higher-level tasks** like genre classification and song segmentation.

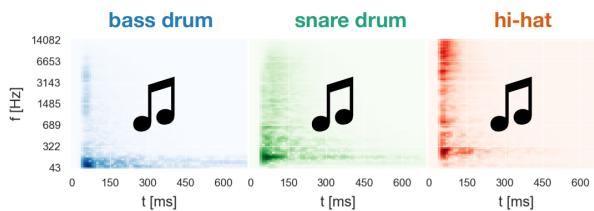


Figure 48: Drum examples



Figure 49: Drum instruments

7.1.2 What are the differences between drum transcription and harmonic instrument transcription?

Drum transcription is generally easier because:

- Drum onsets are well-defined and often broadband.
- There is no need for pitch (f_0) estimation or offset detection.
- The task focuses primarily on timing and instrument classification.

Challenges include:

- Higher temporal accuracy (10–20 ms) requirements.
- Instrument classification is difficult due to spectral overlap.
- Simultaneous onsets and masking effects complicate detection.

Harmonic instruments, on the other hand, produce sounds with distinct pitches and harmonic structures, such as guitars, pianos, and violins. Harmonic transcription involves:

- Estimating the fundamental frequency (f_0) and its harmonics.
- Detecting both onset and offset for note events.
- Handling complex polyphony (multiple overlapping notes).

This makes harmonic transcription more challenging due to the need for precise pitch and temporal accuracy in polyphonic contexts.

7.1.3 What kind of data is needed for automatic drum transcription (training and evaluation)?

ADT requires:

- Audio drum tracks (real recordings, synthesized, or sampled).
- Annotations with high temporal accuracy (e.g., ± 10 ms), specifying drum instrument types and playing techniques.
- Accompaniment tracks or isolated drum tracks for better generalization in machine learning models.

A simpler approach involves using **synthesized audio** derived from symbolic data, such as MIDI files. This method allows for:

- Using MIDI song libraries to synthesize drum and accompaniment tracks separately.
- Extracting annotations directly from MIDI files for drum instrument types and timing.
- Leveraging diverse drum samplers, synthesizers, and other instruments to create varied datasets.

However, synthesized data may not fully represent real-world scenarios. While suitable for evaluation, machine learning models trained on such data must be tested for their ability to generalize to real recordings.

7.1.4 Which metrics are used for ADT evaluation?

Metrics include:

- **Precision:** Proportion of correctly detected drum events.
- **Recall:** Proportion of annotated drum events detected.
- **F1-score:** Harmonic mean of precision and recall, calculated as:

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

7.2 Methods Overview

7.2.1 What are the three categories of ADT methods we discussed in the lecture?

1. **Segment and classify:** Detect onsets and classify them into drum instrument categories.
2. **Separate and detect:** Separate audio sources corresponding to different drum instruments and detect onsets within these sources.
3. **Activation function-based:** Use a single end-to-end system to extract onset detection functions for each instrument directly from the spectrogram.

7.2.2 Explain the segment and classify approach for ADT!

This approach involves two steps

- Detect onsets using basic methods (e.g., energy-based detection).
- Classify each onset into a drum instrument category using rule-based methods (e.g., filters) or machine learning classifiers like support vector machines (SVMs).

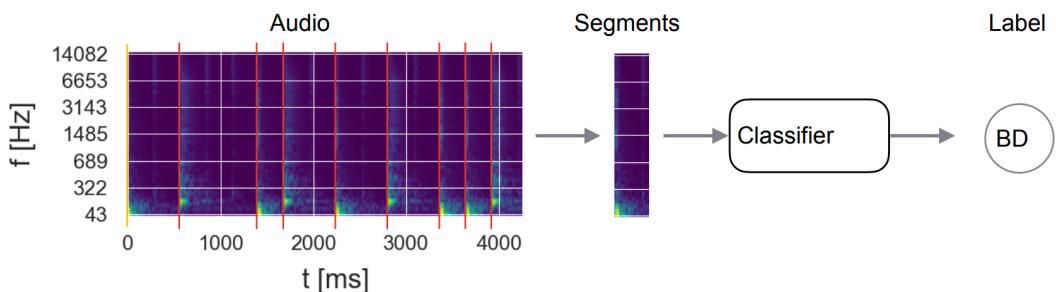


Figure 50: Segment and classify

7.2.3 Explain the separate and detect approach for ADT!

This approach separates audio sources corresponding to different drum instruments before onset detection. Techniques include:

- Bandpass filtering for simple separation.
- Advanced source separation methods like Independent Component Analysis (ICA) or Non-negative Matrix Factorization (NMF).

7.2.4 Explain the activation function based approach for ADT!

The **activation function-based approach** represents the state-of-the-art in automatic drum transcription (ADT). Unlike traditional methods, it directly maps spectrogram inputs to activation functions for each drum instrument in a single processing step. Peaks in the activation functions correspond to drum onsets, streamlining the transcription process and improving accuracy.

Challenges in Previous Methods:

- **Segment and Classify:** This method involves first detecting onsets and then classifying each detected onset into instrument categories. Challenges include:

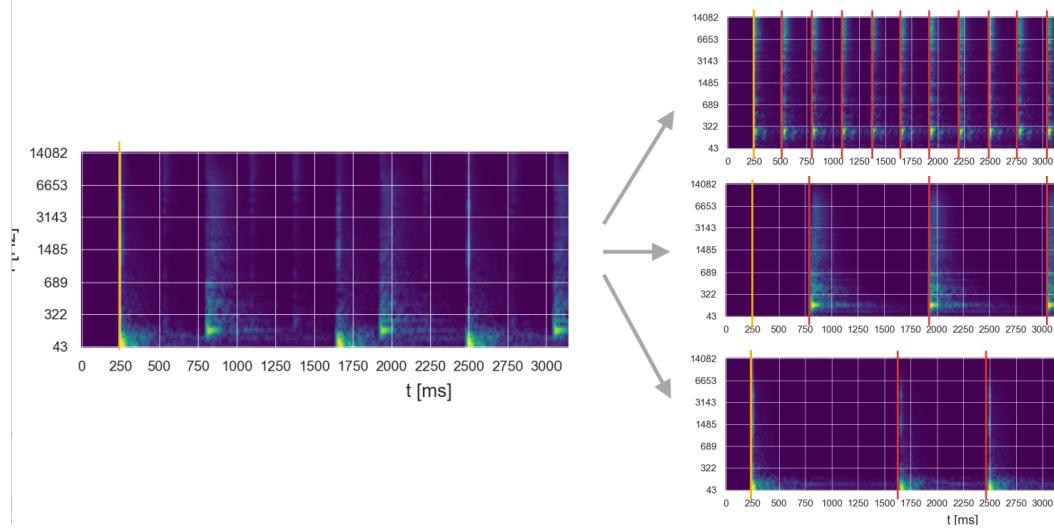


Figure 51: Separate and select

- High sensitivity to onset detection accuracy.
- Limited capability to handle overlapping or simultaneous drum hits.

- **Separate and Detect:** This method separates the audio into individual instrument sources before detecting onsets for each source. Challenges include:
 - Ineffectiveness of simple bandpass filtering for complex signals.
 - Overfitting and parameterization difficulties in advanced source separation techniques.

How Activation Function-Based Approaches Overcome These Challenges:

- By treating the transcription task as a single step, end-to-end methods reduce the reliance on intermediate processes that can introduce errors.
- Neural networks and NMF-based systems can directly learn to recognize patterns and temporal structures in the spectrogram, effectively addressing issues with overlapping onsets and spectral masking.

The two primary implementations of this approach are:

- **NMF-Based Approaches:** These use matrix decomposition techniques to identify activation patterns for each instrument.
- **Neural Network-Based Approaches:** These leverage architectures like CNNs, RNNs, and CRNNs to learn complex relationships directly from spectrograms.

7.3 Non-Negative Matrix Factorization

7.3.1 Explain in your own words the algorithm for non-negative matrix factorization!

Non-negative Matrix Factorization (NMF) approximates a spectrogram X as the product of two matrices:

- B : Basis matrix representing spectral templates for instruments.

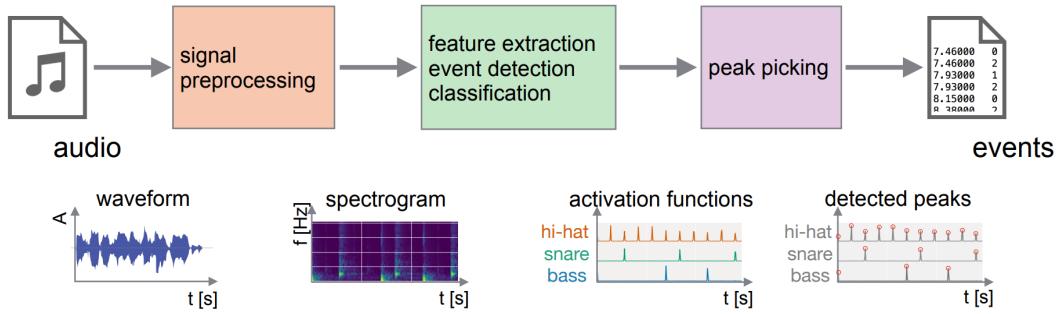


Figure 52: Activation Function Based

- G : Activation matrix indicating the presence of each template over time.

The algorithm iteratively updates B and G using multiplicative update rules to minimize a loss function, typically the Kullback-Leibler divergence.

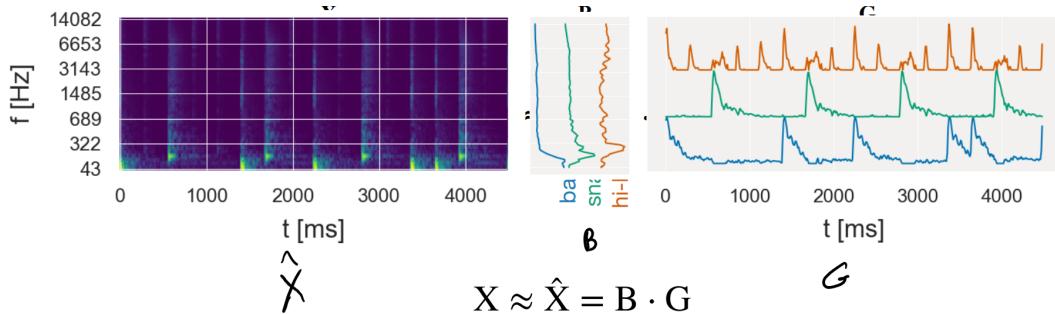


Figure 53: NMF basic idea

7.3.2 What are the challenges that come with NMF for automatic drum transcription?

Challenges include:

- Defining the number of templates (size of B).
- Adapting to varying spectral energy distributions over time.
- Handling harmonic and accompaniment content, which can overlap with drum components.
- Assigning templates to specific instruments for meaningful transcription.

7.3.3 Name variants for NMF in the context of ADT, what are the differences?

In the context of NMF, a **template** refers to a spectral representation of an instrument's sound, typically derived from the spectrogram of individual drum hits. Templates serve as the basis elements in the decomposition process, representing the characteristic frequency content of specific instruments. Below are the main NMF variants:

- **Fixed templates:**

- Templates are predefined using spectrograms of individual drum hits.

- The mean spectrum of each drum type (e.g., bass drum, snare drum) is used as the base.
- The basis matrix B is fixed and not updated during the decomposition process.
- This approach works well for solo tracks with clean, isolated hits but struggles to adapt to different sounding instruments or variations in timbre.

- **Semi-adaptive NMF:**

- Combines fixed templates with adaptive updates to the basis matrix B during later iterations.
- Initially uses predefined templates, blending them with updates to allow some flexibility.
- Provides a balance between specificity and adaptability, making it suitable for more diverse datasets.

- **Partially-fixed NMF:**

- Adds additional adaptive templates to model harmonic or accompaniment content alongside drum sounds.
- Updates the drum templates (e.g., bass drum) while keeping the harmonic templates fixed during most iterations.
- Adjusts the basis matrix B and activation matrix G iteratively:
 - * Harmonic templates (B_H) and activations (G_H) are updated based on the harmonic portion of the signal (\hat{X}_H).
 - * Drum templates (B_D) and activations (G_D) are updated for drum signals (\hat{X}_D), ensuring separation between drum and harmonic components.
- This approach effectively models accompaniment while maintaining high-quality drum transcription.

- **Convolutional NMF:**

- Uses 2D templates that account for both spectral and temporal variations in drum sounds.
- Models the time-evolving nature of drum hits, capturing changes in spectral energy distribution over time.
- Offers improved transcription accuracy but at the cost of higher computational complexity.

7.4 NN Based Methods

7.4.1 Explain the setup for neural network based ADT!

The setup involves:

- Preprocessing audio to generate spectrograms.
- Training a neural network (e.g., CNN, RNN, or TCN) with annotated drum events to predict activation functions for each instrument.
- Post-processing activation functions to identify drum onsets and classify events.

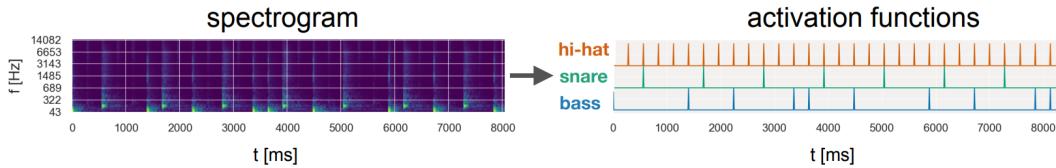


Figure 54: NN based method example

7.4.2 Explain the difference between the multi-task and single-task NN approach for ADT!

- **Single-task NN:** A separate network is trained for each drum instrument. While effective, this approach requires more training time as the number of instruments increases, making it less efficient for complex datasets.
- **Multi-task NN:** A single network with multiple output layers predicts activation functions for all instruments simultaneously. This approach:
 - Exploits shared features across instruments, reducing redundancy and improving efficiency.
 - Can include outputs for beat and downbeat tracking, leveraging meta-data like beat time to enhance transcription performance.
 - Provides better performance in complex scenarios, such as with accompaniment or masked onsets (e.g., overlapping hi-hat sounds).
 - Requires a sufficiently large receptive field to capture context, allowing the network to model relationships across time effectively.

7.4.3 What are questions to answer when trying to choose a model for a given task?

- Can the architecture model the problem efficiently?
- Does it have a sufficiently large receptive field for context-aware predictions?
- Is it computationally feasible with acceptable model size?
- Can it be trained efficiently on the available data?

7.4.4 Explain the differences between dense, recurrent, convolutional, and dilated-conv neural networks! What are the advantages and disadvantages of each architecture category?

- **Dense NN:** Fully connected layers.
 - **Advantages:** Simple and effective for basic tasks.
 - **Disadvantages:** Unsuitable for time-series data, lacks temporal context.
- **Recurrent NN (RNN):** Captures temporal dependencies in sequential data.
 - **Advantages:** Effective for modeling sequences and temporal patterns.
 - **Disadvantages:** Hard to train due to vanishing gradients, computationally expensive.
- **Convolutional NN (CNN):** Extracts local features from spectrograms.
 - **Advantages:** Efficient for spatial feature extraction.
 - **Disadvantages:** Limited temporal context unless stacked deeply.
- **Dilated CNN (TCN):** Uses dilated convolutions for modeling long-range dependencies.
 - **Advantages:** Provides a large receptive field with fewer layers, balancing efficiency and context modeling.
 - **Disadvantages:** Can be more complex to implement compared to standard CNNs.

7.4.5 Explain how a CNN/RNN/TCN can be used to perform automatic drum transcription! How is the spectrogram processed in the context for each model?

- **CNN:** Processes spectrograms frame-by-frame, focusing on spatial features that represent drum events. CNN architectures, such as VGG-style CNNs, include layers of convolution, batch normalization, and pooling to extract meaningful patterns. For example, a CNN might use two sets of convolutional layers followed by max-pooling to reduce dimensions and capture local features efficiently.
- **RNN:** Processes sequential frames to capture temporal dependencies in the spectrogram. Bidirectional RNNs with GRUs are commonly used, as they consider both past and future contexts while predicting drum onsets. These architectures are well-suited for time-series tasks, but their training can be computationally intensive.
- **CRNN (CNN + RNN):** Combines CNNs for acoustic modeling with RNNs for music language modeling. The CNN extracts features from the spectrogram, which are then passed to the RNN for sequence modeling, enabling better performance on tasks requiring both spatial and temporal understanding.
- **TCN (Temporal Convolutional Network):** Uses dilated convolutions to capture long-range dependencies in the spectrogram. TCN architectures typically include stacks of dilated convolutional layers to expand the receptive field while maintaining computational efficiency. These are well-suited for long sequences and are easier to train than RNNs.

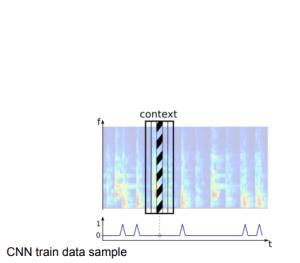


Figure 55: CNN architecture

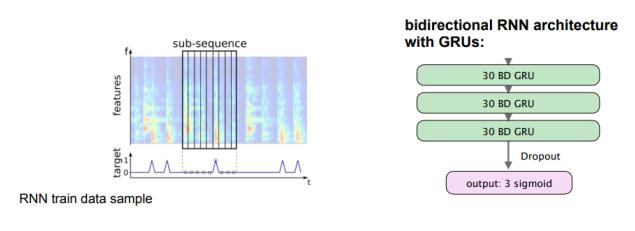


Figure 56: RNN architecture

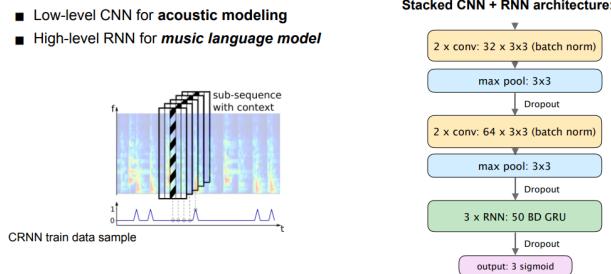


Figure 57: CRNN architecture

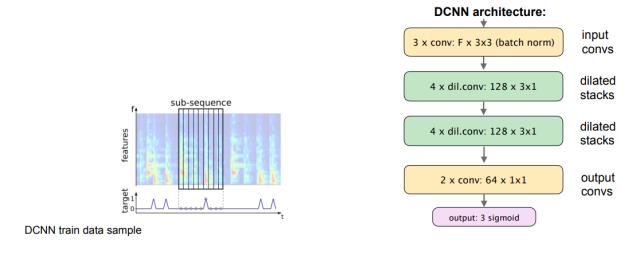


Figure 58: DCNN (TCN) architecture

8 Audio Representation Learning

8.1 Introduction to Audio Representation Learning

Representation learning is a fundamental aspect of machine learning. The effectiveness of machine learning algorithms is highly dependent on the quality of data representation. Traditionally, feature engineering has been labor-intensive, requiring significant domain expertise. Modern representation learning aims to automate this process, extracting meaningful representations directly from raw data, thus bridging the gap between raw signals and machine learning models.

8.2 Machine Learning Paradigms

Machine learning paradigms have evolved significantly, with representation learning and deep learning becoming the focus of modern approaches. Figure ?? illustrates the four main paradigms of machine learning.

8.2.1 Rule-Based Systems

- Algorithms are written manually with explicit instructions for solving a task.
- Strength: Effective for well-defined problems.
- Weakness: Inflexible and requires significant human intervention.

8.2.2 Classic Machine Learning

- Feature extractors are designed manually, and a classifier is trained on top of these features.
- Strength: Reduces the need for rule-writing.

- Weakness: Relies heavily on human expertise for feature engineering.

8.2.3 Representation Learning

- Feature extractors are learned automatically, often using unsupervised techniques, and a classifier is trained on the learned representations.
- Strength: Reduces the need for manual feature engineering and improves adaptability.
- Weakness: Requires sufficient data and computational resources to learn meaningful representations.

8.2.4 Deep Learning

- Deep learning stacks multiple layers of simpler functions, often in the form of neural networks, to map inputs to outputs.
- Strength: Can process raw data end-to-end, optimizing both feature extraction and classification simultaneously.
- Weakness: Computationally expensive and requires large labeled datasets.

8.3 Audio Feature Extraction

Audio representation involves extracting features at different levels:

- **Low-level features:** Pitch, timbre, and rhythm.
- **Mid-level features:** Genres, styles, and instruments.
- **High-level features:** Themes, emotions, and contextual metadata.

Key features include:

- Mel-frequency cepstral coefficients (MFCCs).
- Chroma features.
- Rhythm patterns and statistical spectrum descriptors.
- Spectral envelope characteristics such as centroid, contrast, and rolloff.

8.4 Signal Processing for Audio Representations

Audio representations rely on transformations from the time domain to the frequency domain:

- Time-domain analysis (waveforms).
- Frequency-domain analysis (spectra) using Fourier Transform (FFT), Discrete Cosine Transform (DCT), and Wavelet Transform.
- Time-frequency representations, such as spectrograms.

Figure ?? illustrates the progression from raw audio to advanced feature extraction techniques.

8.5 Approaches to Audio Representation Learning

8.5.1 Linear Methods

- Principal Component Analysis (PCA).
- Linear Discriminant Analysis (LDA).

8.5.2 Nonlinear Methods

- Kernel PCA.
- Restricted Boltzmann Machines (RBMs).
- Manifold learning techniques.

8.5.3 Deep Representation Learning

Deep learning has revolutionized representation learning with methods such as:

- **Autoencoders (AE):** Learn compressed latent representations of data by encoding and decoding.
- **Triplet Networks:** Use triplet loss to learn embeddings based on similarity.
- **Supervised Learning:** Train neural networks for specific tasks and extract latent representations.

8.6 Deep Learning Methods for Audio

8.6.1 Autoencoders (AEs)

Autoencoders compress input data into a latent space and reconstruct it. The latent representation encodes essential information while discarding noise. Variants include:

- **Undercomplete AEs:** Force compression by using a smaller latent space than the input.
- **Overcomplete AEs:** Use a larger latent space but require regularization to prevent trivial solutions.

8.6.2 Supervised Learning for Representation Extraction

In supervised learning:

- Train a deep neural network for classification or regression tasks.
- Use the latent representation from a hidden layer for downstream tasks.

8.6.3 Triplet Networks

Triplet networks optimize embeddings by minimizing the distance between an anchor and a positive sample while maximizing the distance to a negative sample. The triplet loss function ensures meaningful separation in the latent space.

$$L = \max\{d(a, p) - d(a, n) + m, 0\}$$

where d is the distance metric, and m is the margin.

8.7 Challenges in Audio Representation Learning

- **Task-specific Representations:** Representations often fail to generalize across tasks.
- **Semantic Gap:** Bridging the gap between raw audio features and high-level semantic understanding.
- **Data Annotation:** Annotating large datasets is expensive and time-consuming.
- **Contextual Variance:** Variations in recording quality, styles, and genres pose challenges.

8.8 Applications of Audio Representation Learning

- **Keyword Spotting:** Few-shot learning methods for detecting specific phrases.
- **Speaker Diarization:** Differentiating speakers in multi-speaker audio.
- **Music Representation:** Tasks such as genre classification, similarity retrieval, and recommendation systems.

9 Acoustic Scene Classification

9.1 Introduction to Acoustic Scene Classification (ASC)

Acoustic Scene Classification (ASC) involves identifying and labeling the acoustic environment of an audio stream based on predefined categories. This task plays a critical role in various applications, from security and urban planning to multimedia analysis.

9.2 Scene Classification vs. Sound Event Detection

- **Acoustic Scene Classification (ASC):**

- Single-label classification.
- No onset detection.
- Labels apply to the entire track.

- **Sound Event Detection (SED):**

- Multi-label classification.
- With or without onset detection.
- Focus on detecting and categorizing specific sound events, including their temporal boundaries.

9.3 Applications of ASC and SED

9.3.1 Health and Medical Sector

- Telemonitoring of patients for fall detection and speech/non-speech classification.
- Ambient Assisted Living (AAL):
 - Anticipate inhabitants' needs unobtrusively.
 - Smart home controls: e.g., turning on monitors when a baby cries.

9.3.2 Security and Surveillance

- Detecting dangerous sounds in public places, tunnels, and elevators.
- Forensic media analysis after crimes or attacks, identifying gunshots or explosions.

9.3.3 Ecosystem Monitoring

- Identification of animal sounds to estimate population densities, migration patterns, and habitat health.

9.3.4 Other Applications

- Urban soundscape analysis for noise pollution classification.
- Industrial applications, such as monitoring production lines, machinery, and motors.

9.4 Challenges in Acoustic Scene Classification and Sound Event Detection

9.4.1 Generalization Challenges

- **Sensor Variability:** Different microphones have unique frequency responses and directionality characteristics.
- **Location Variability:** Indoor versus outdoor environments exhibit different acoustic properties.
- **Temporal Variations:** Day versus night, seasonal changes, and weather conditions.

9.4.2 Data Augmentation to Address Challenges

Data augmentation techniques improve model robustness by enhancing training datasets:

- Noise addition.
- Pitch shifting and time stretching.
- Mixing multiple audio recordings or using block mixing.
- Temporal and spectral masking (e.g., cut-out techniques).

9.5 Data for ASC and SED

9.5.1 Ground Truth Label Types

- **Strong Labels:** Precise annotations with onset and offset times, typically at a high cost.
- **Weak Labels:** Coarse annotations at the clip level, enabling larger datasets.

9.5.2 Data Annotation Techniques

Crowdsourcing tools and platforms facilitate annotation, balancing accuracy with scalability:

- Examples: Audio-Annotator, Label Studio, and Wavesurfer.
- Strategies:
 - Use spectrogram visualizations for annotators.
 - Employ training phases to improve annotation quality.

9.6 Approaches to ASC and SED

9.6.1 Machine Learning for Acoustic Scene Classification

- **Bag-of-Frames Approach:** Statistical distributions of local spectral features, such as Mel-frequency Cepstral Coefficients (MFCCs).
- **High-Level Feature Extraction:** Non-negative Matrix Factorization (NMF) to derive acoustic bases for classification.

9.6.2 Deep Learning for ASC and SED

- **Convolutional Neural Networks (CNNs):** Learn spectral patterns from input spectrograms.
- **Recurrent Neural Networks (RNNs):** Capture temporal dependencies in audio signals.
- **CRNNs (Convolutional-Recurrent Neural Networks):** Combine CNNs and RNNs for end-to-end learning of spectral and temporal features.
- **Attention Mechanisms:** Enhance focus on relevant audio regions, improving convergence and generalization.

9.7 Evaluation Campaigns and Datasets

9.7.1 Key Evaluation Campaigns

- **DCASE:** Detection and Classification of Acoustic Scenes and Events. Annual competition with datasets and benchmarks.
- **LifeClef:** Focused on bioacoustics, such as bird song identification.

9.7.2 Popular Datasets

- **Google Audio Set:** Over 2M videos annotated with 632 audio event categories.
- **DESED:** Domestic environment sounds with both labeled and unlabeled data.